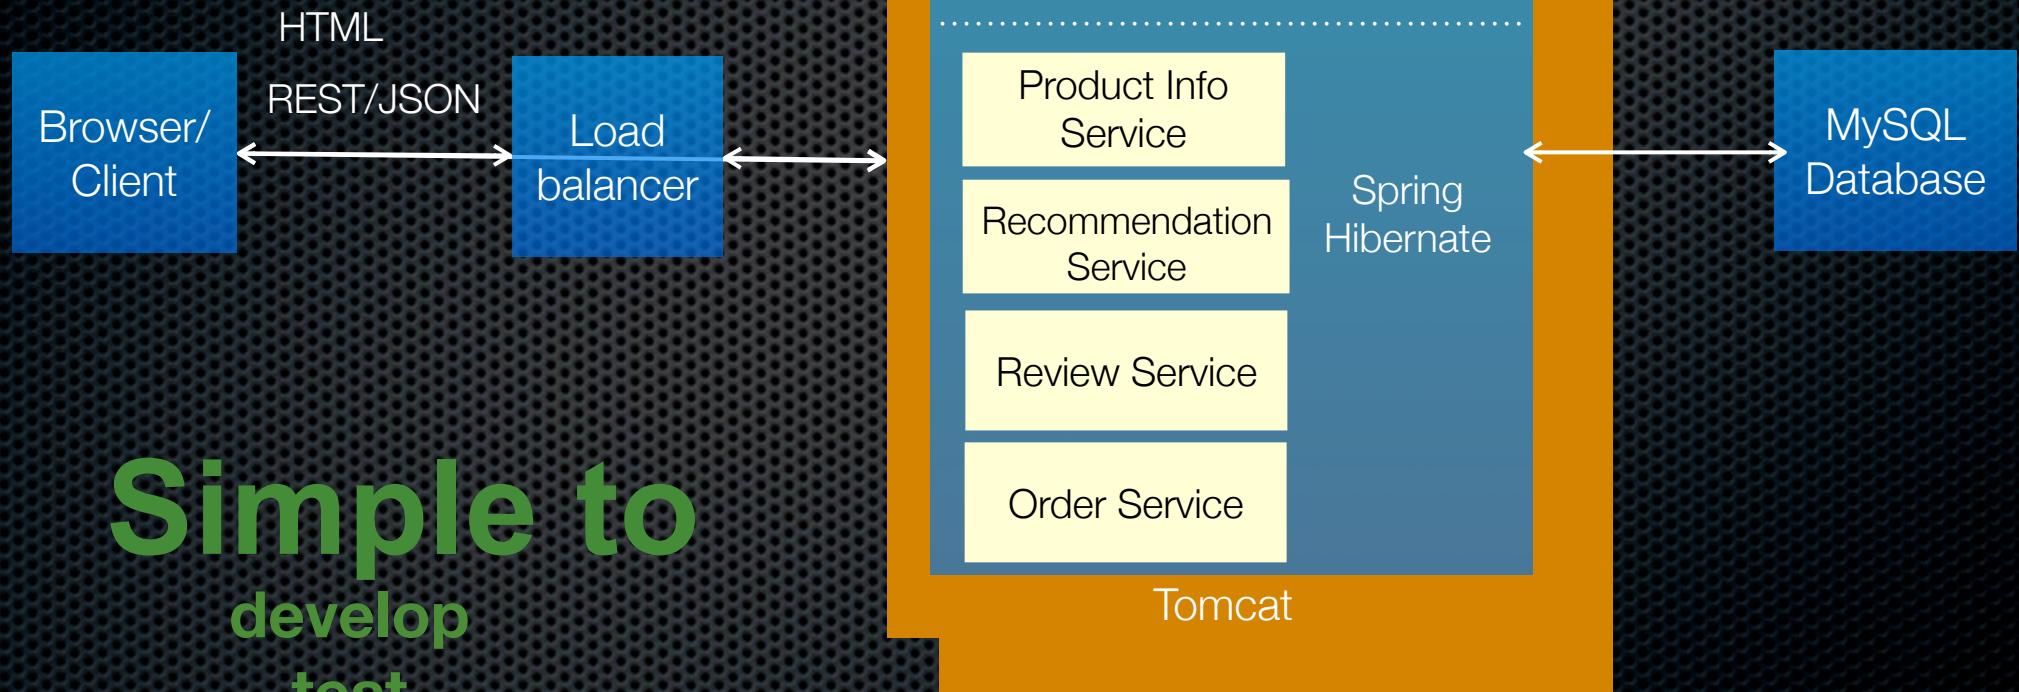


Let's imagine you are
building an online store

Traditional application architecture



Simple to
develop
test
deploy
scale

But large, complex, monolithic
applications



problems

Intimidates developers



Obstacle to frequent deployments

- Need to redeploy everything to change one component
- Interrupts long running background (e.g. Quartz) jobs
- Increases risk of failure

Eggs in
one basket



- Updates will happen less often - really long QA cycles
- e.g. Makes A/B testing UI really difficult

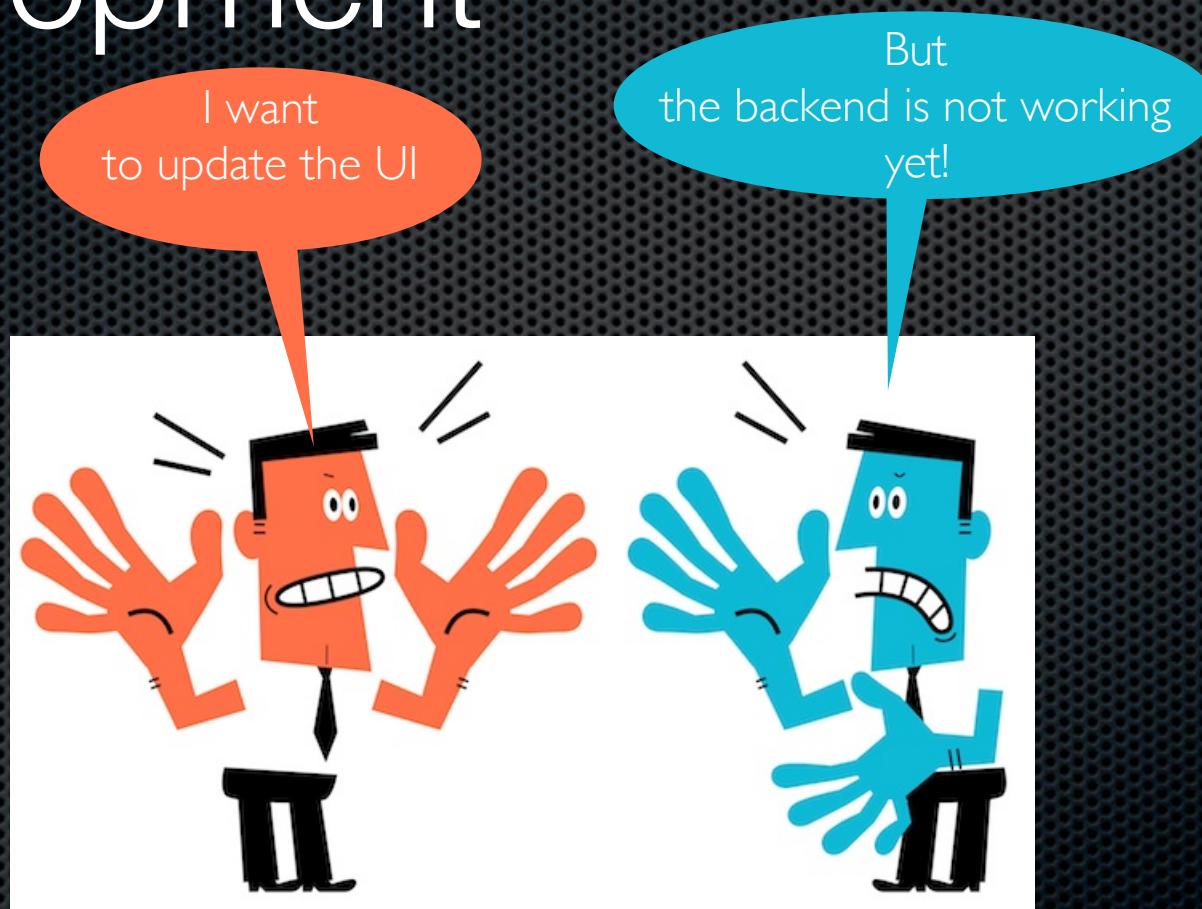
Overloads your IDE and container



Slows down development

@crichtson

Obstacle to scaling development



Lots of coordination and communication required

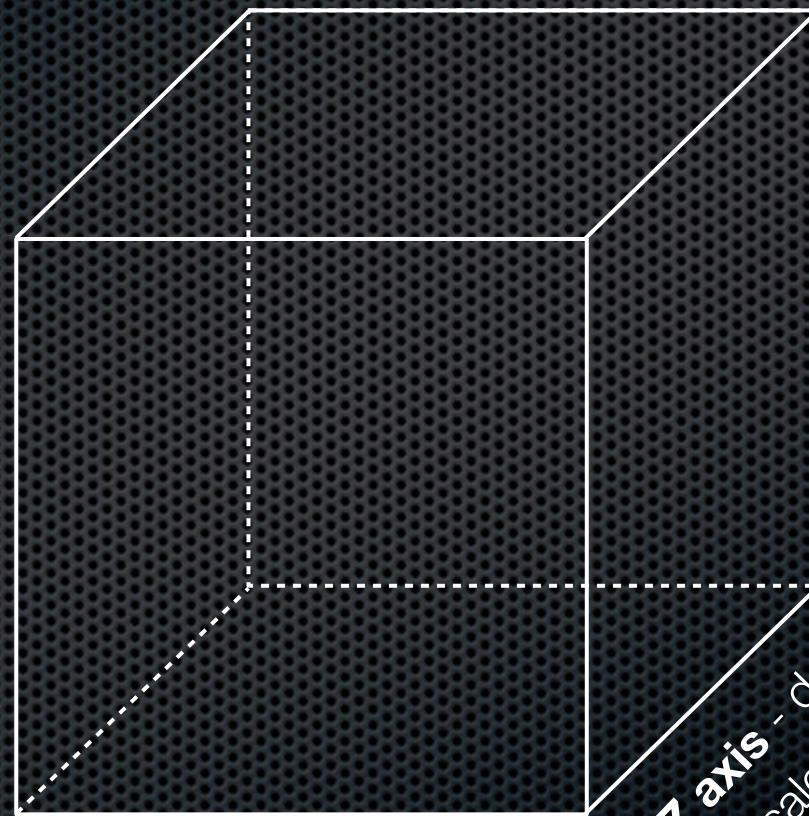
Requires long-term commitment to a technology stack



The scale cube

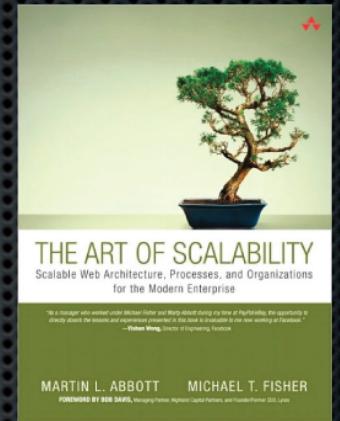
Y axis -
functional
decomposition

Scale by
splitting
different things



X axis
- horizontal duplication

Z axis - data partitioning
Scale by splitting similar
things

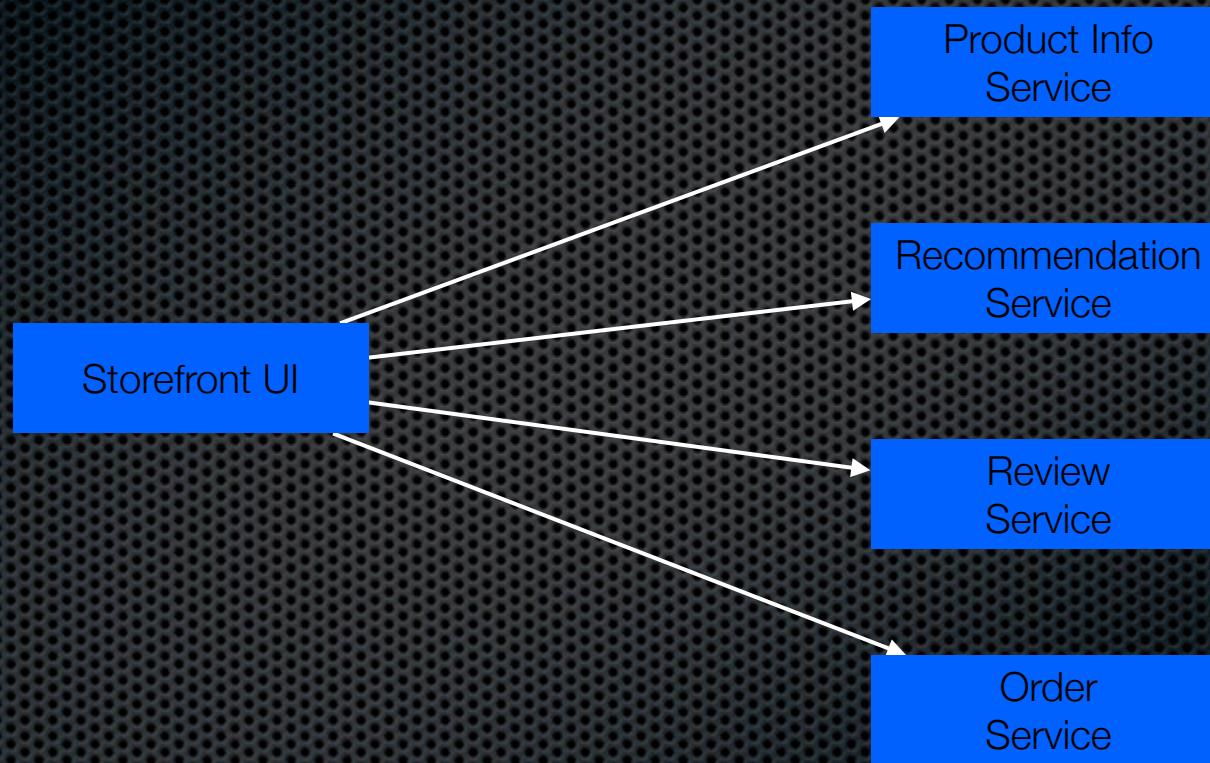


@crichardson

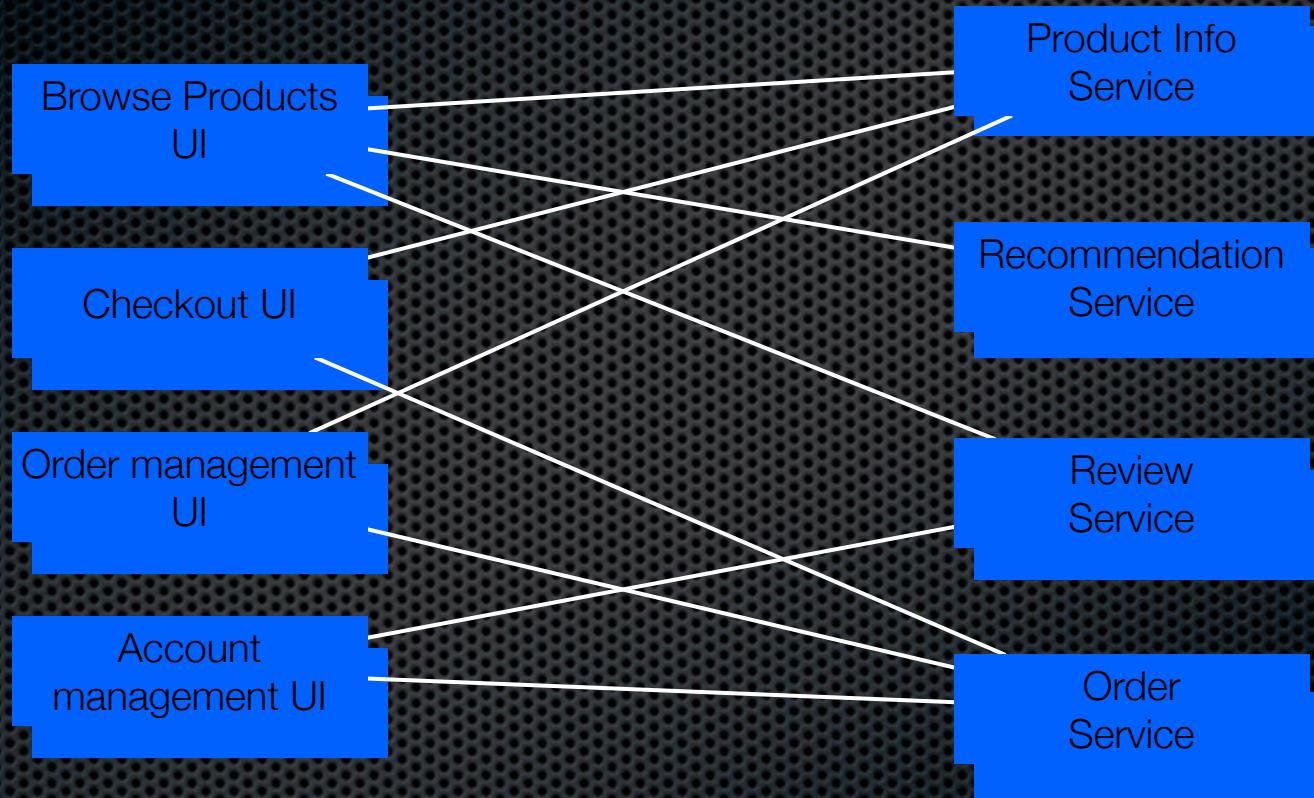
Y-axis scaling - application level



Y-axis scaling - application level



Y-axis scaling - application level



Apply X-axis and Z-axis scaling
to each service independently

Service deployment options

Isolation, manageability

- VM or Physical Machine
- Docker/Linux container
- JVM
- JAR/WAR/OSGI bundle/...



Density/efficiency

Partitioning strategies...

- ❖ Partition by noun, e.g. product info service
- ❖ Partition by verb, e.g. Checkout UI
- ❖ Single Responsibility Principle
- ❖ Unix utilities - do one focussed thing well

Partitioning strategies

- ❖ Too few
 - ❖ Drawbacks of the monolithic architecture
- ❖ Too many - a.k.a. Nano-service anti-pattern
 - ❖ Runtime overhead
 - ❖ **Potential** risk of excessive network hops
 - ❖ **Potentially** difficult to understand system

Something of an art

More service, less micro

But more realistically...

Focus on building services that make development and deployment easier

- not just tiny services

Real world examples



<http://techblog.netflix.com/>

~600 services



<http://highscalability.com/amazon-architecture>

100-150 services to build a page



<http://www.addsimplicity.com/downloads/eBaySDForum2006-11-29.pdf>

<http://queue.acm.org/detail.cfm?id=1394128>

There are many benefits

Smaller, simpler apps

- Easier to understand and develop
- Less jar/classpath hell - who needs OSGI?
- Faster to build and deploy
- Reduced startup time - important for GAE

Scales development:
develop, deploy and scale
each service independently

Improves fault isolation

Eliminates long-term commitment
to a single technology stack



Modular, polyglot, multi-
framework applications

Two levels of architecture

System-level

Services

Inter-service glue: interfaces and communication mechanisms

Slow changing

Service-level

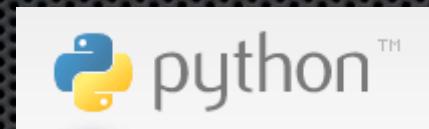
Internal architecture of each service

Each service could use a different technology stack

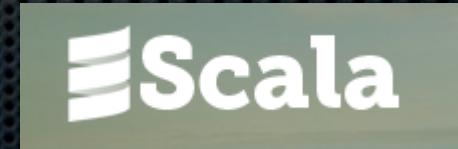
Pick the best tool for the job

Rapidly evolving

Easily try other technologies



Spring Boot



... and fail safely

But there are drawbacks

Complexity

Complexity of developing a distributed system

<http://highscalability.com/blog/2014/4/8/microservices-not-a-free-lunch.html>

Multiple databases & Transaction management

e.g. Fun with eventual consistency

Complexity of testing a distributed system

<http://highscalability.com/blog/2014/4/8/microservices-not-a-free-lunch.html>

Complexity of deploying and operating a distributed system

<http://highscalability.com/blog/2014/4/8/microservices-not-a-free-lunch.html>

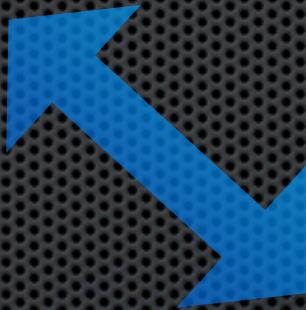
You need a lot of automation

Developing and deploying
features that span multiple
services requires careful
coordination

When to use it?

In the beginning:

- You don't need it
- It will slow you down



Later on:

- You need it
- Refactoring is painful

Let's imagine that you want to display a product's details...

Product Info

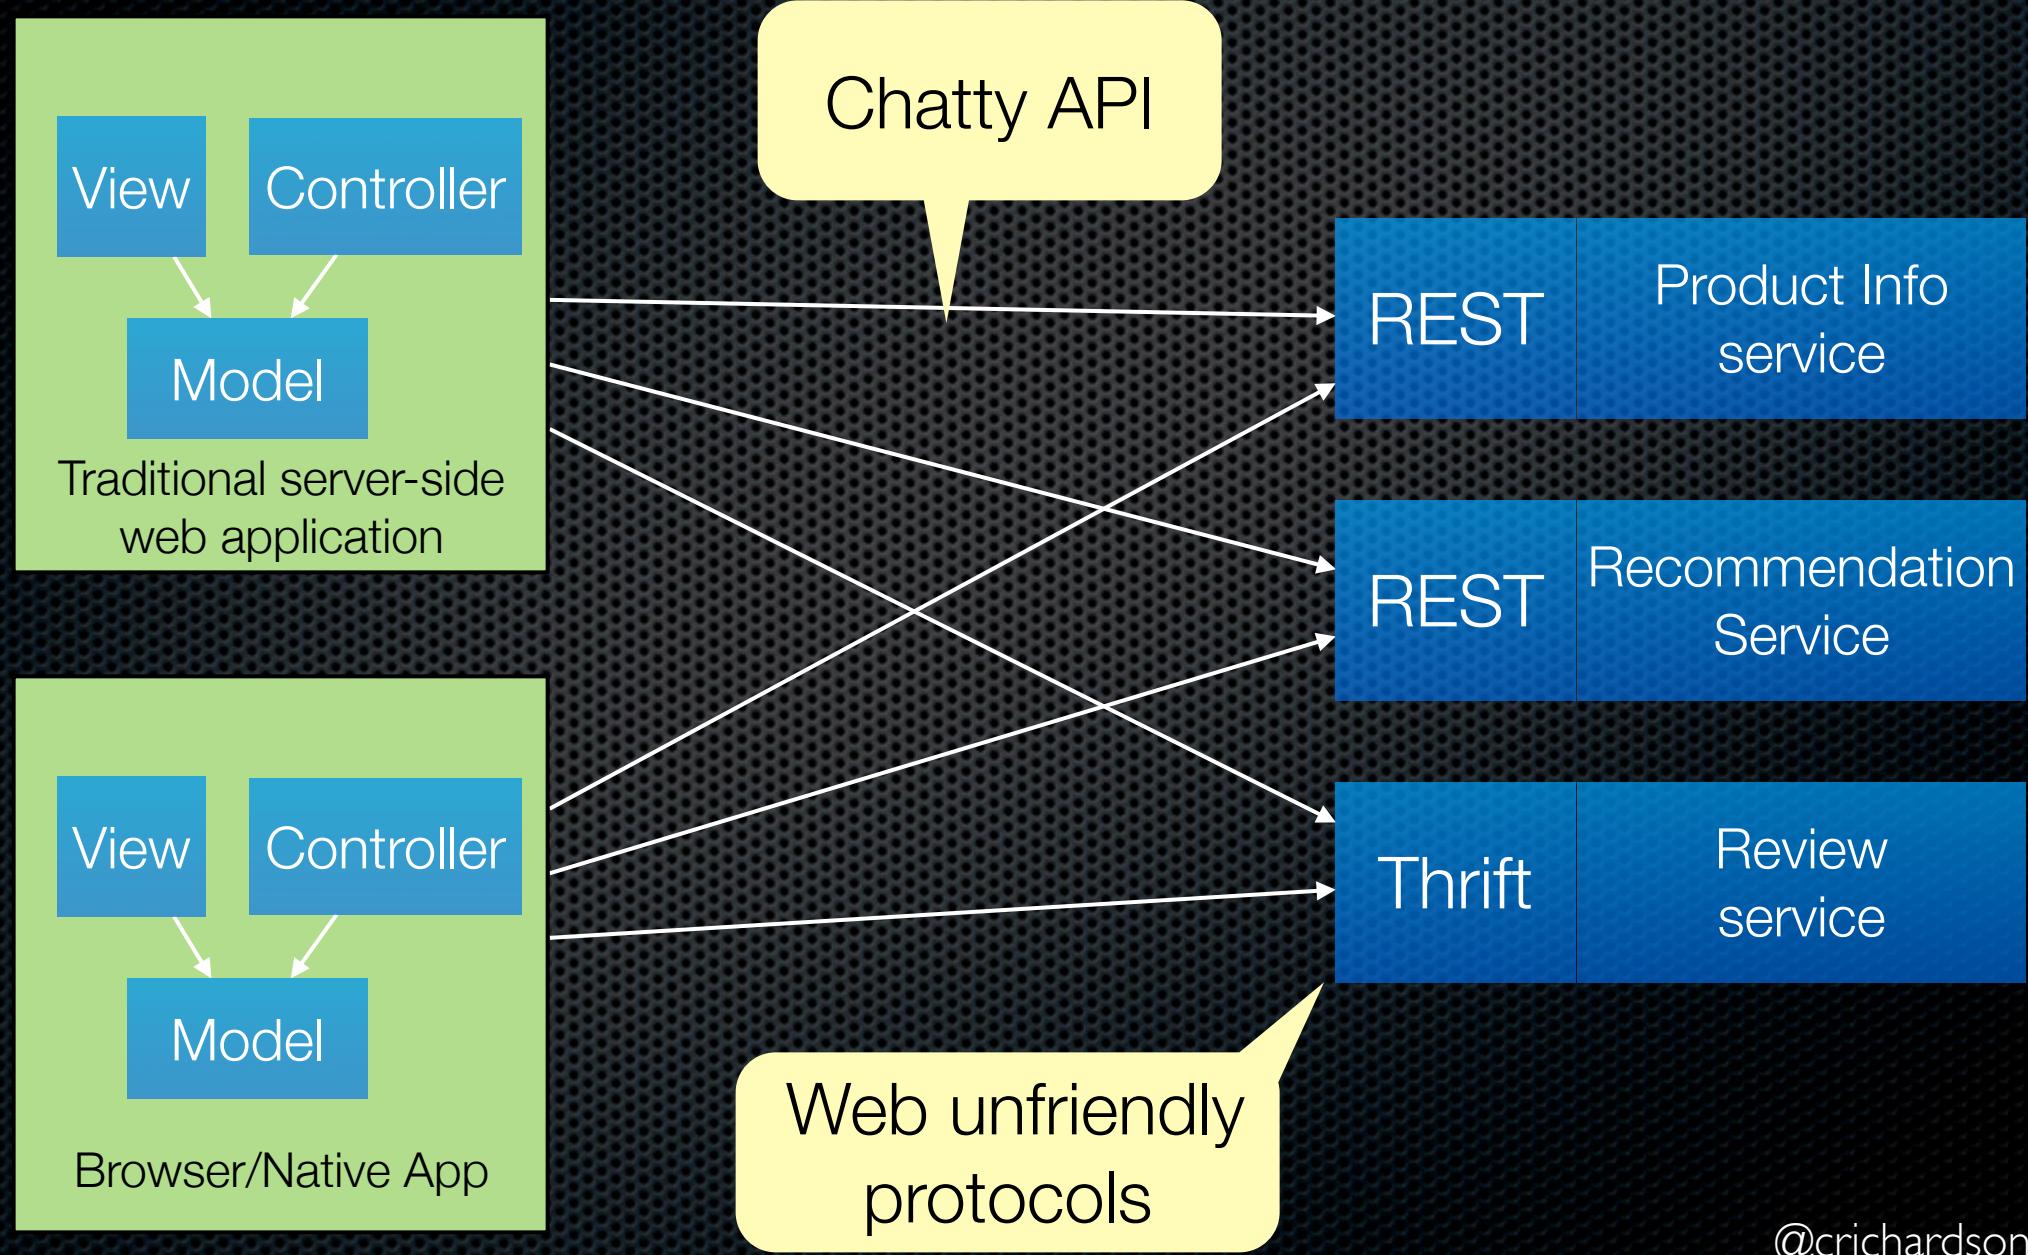
The screenshot shows the product page for 'POJOs in Action: Developing Enterprise Applications with Lightweight Frameworks' by Chris Richardson. The page includes the book cover, price (\$34.53), customer reviews (4.5 stars from 31 reviews), and a 'Buy New' button. Below the main product information, there are sections for 'Frequently Bought Together' and 'Customers Who Bought This Item Also Bought', each listing several related books with their titles, authors, and prices.

Reviews

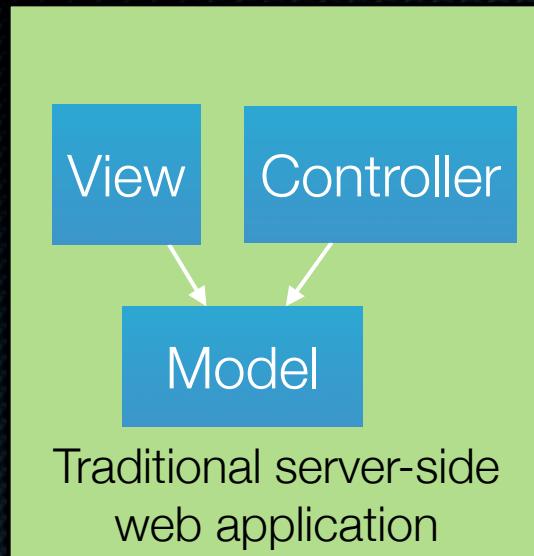
Recommendations

@crichtson

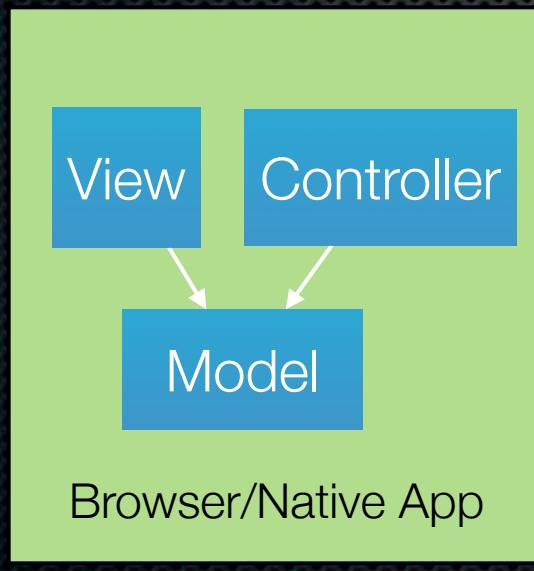
Directly connecting the front-end to the backend



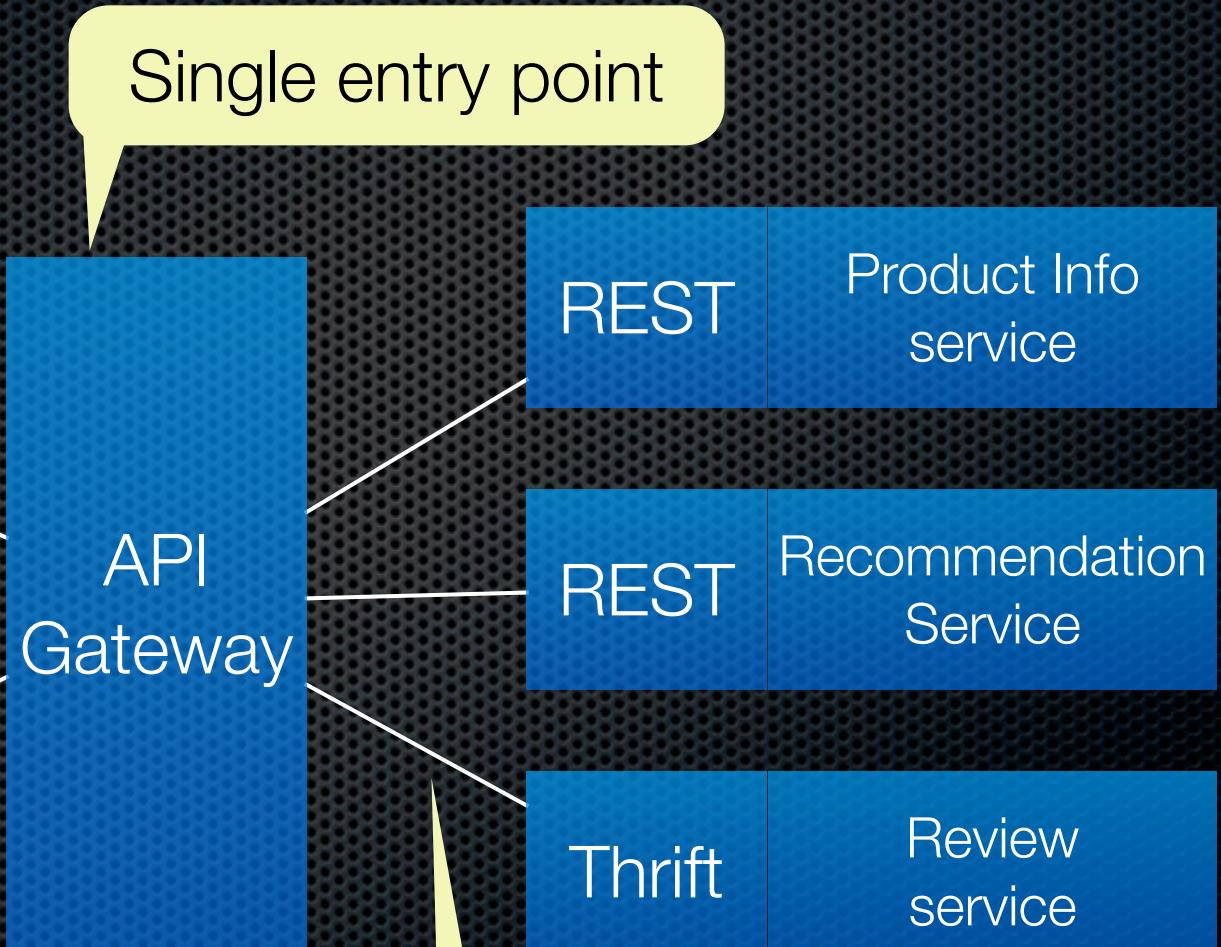
Use an API gateway



Traditional server-side
web application



Browser/Native App



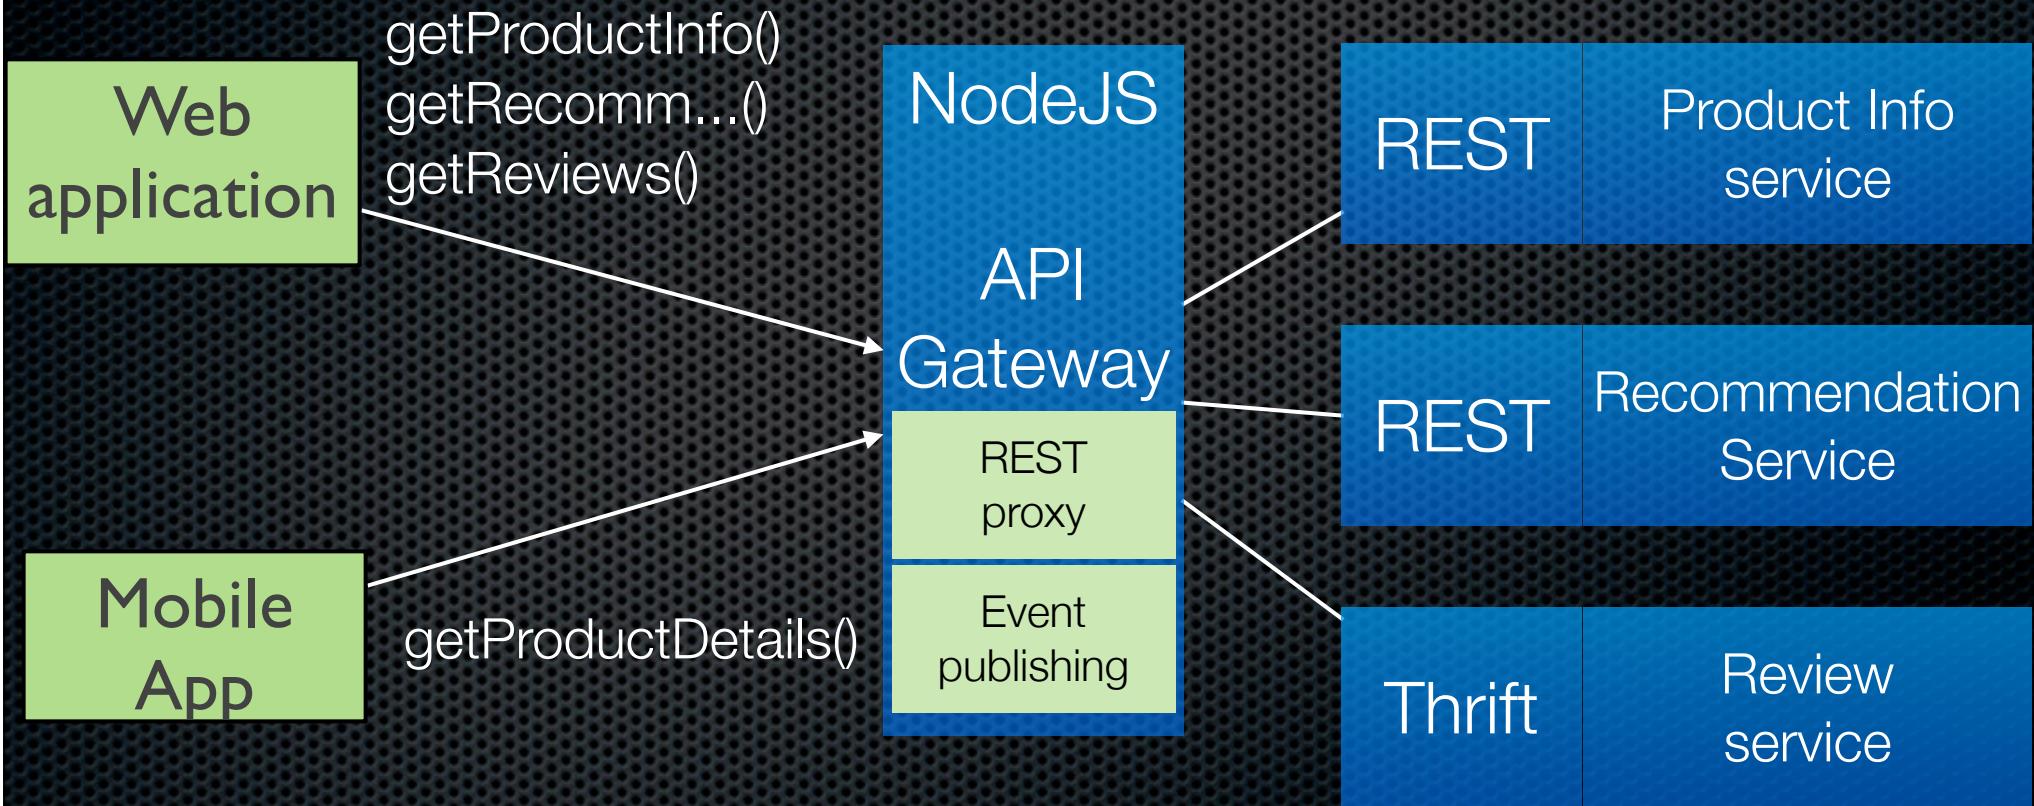
Single entry point

API
Gateway

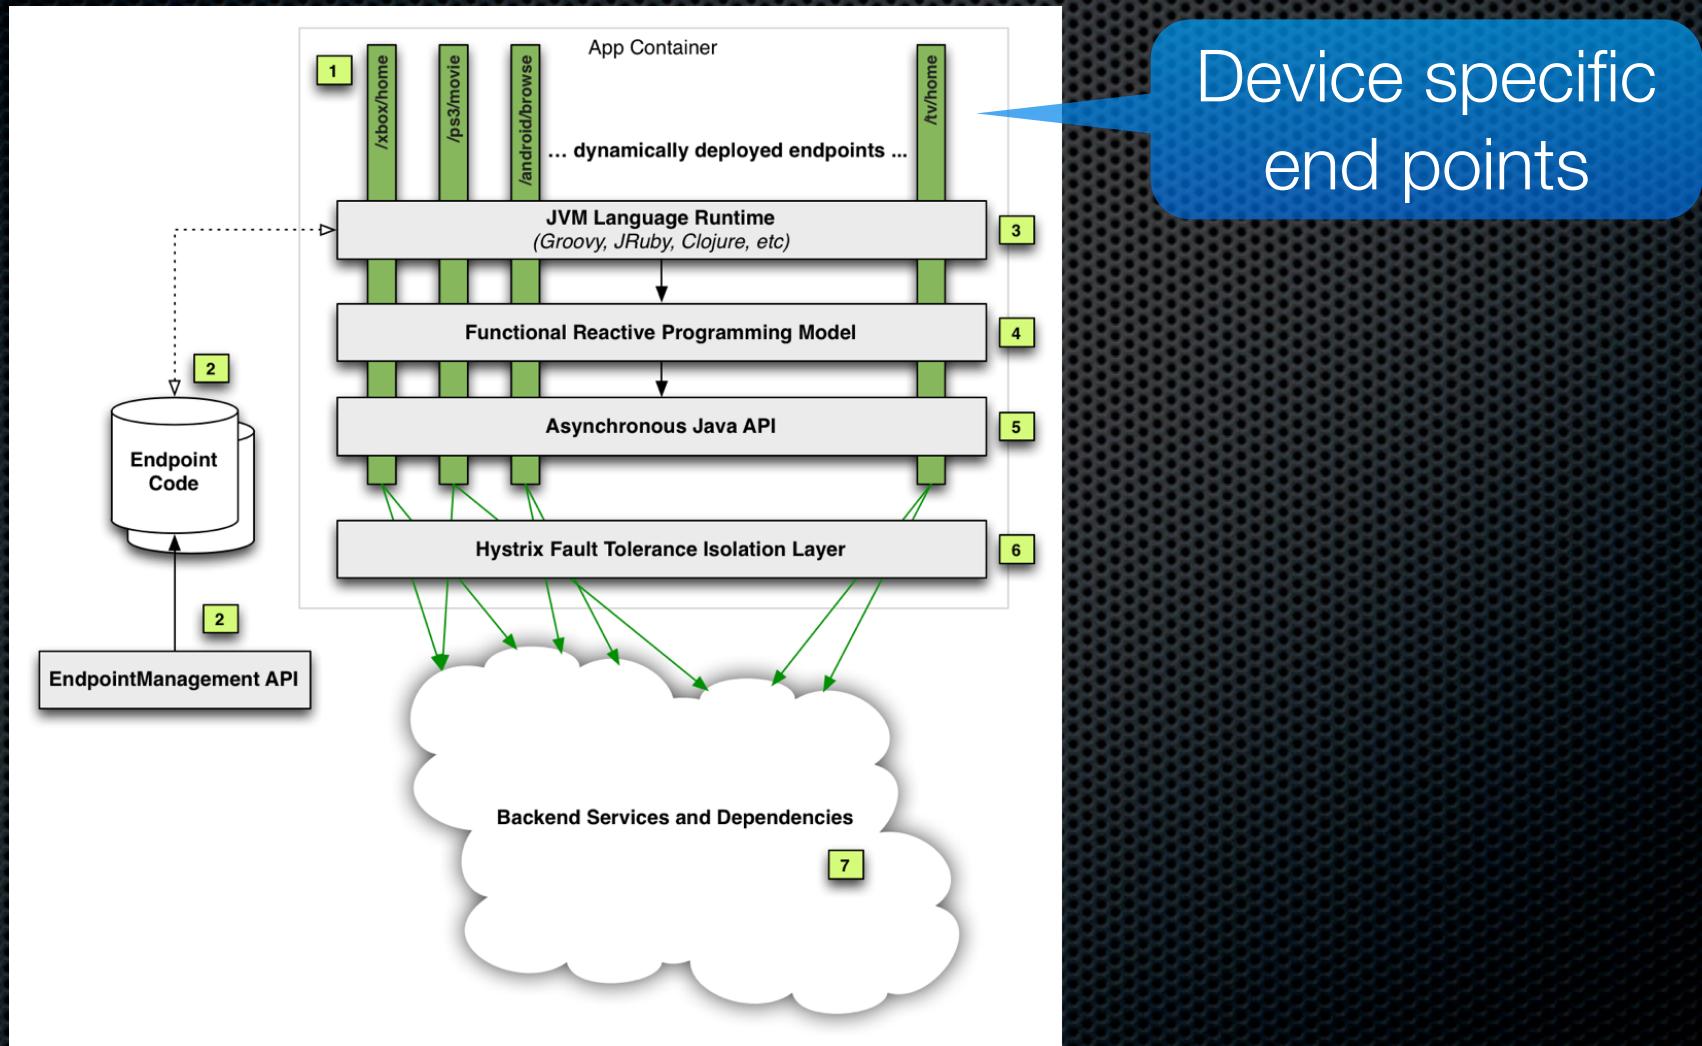
Client
specific APIs

Protocol
translation

Optimized client-specific APIs



Netflix API Gateway



<http://techblog.netflix.com/2013/01/optimizing-netflix-api.html>

@crichton

API gateway design challenges

- ❖ Performance and scalability
 - ❖ Non-blocking I/O
 - ❖ Asynchronous, concurrent code
- ❖ Handling partial failures
- ❖

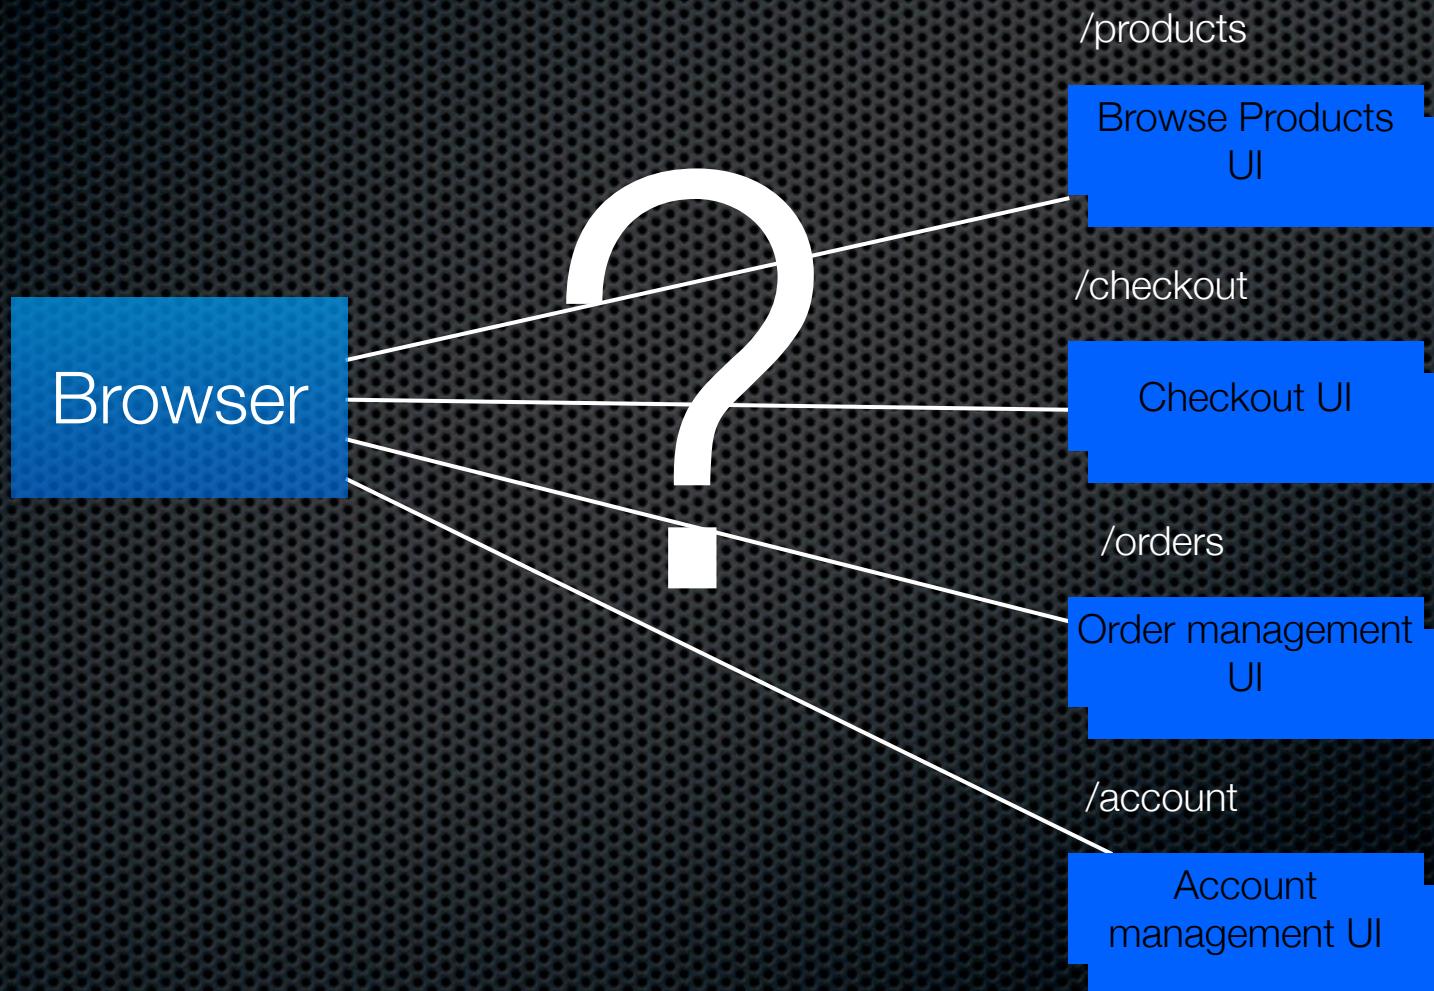
<http://techblog.netflix.com/2012/02/fault-tolerance-in-high-volume.html>

Useful frameworks for building an API gateway

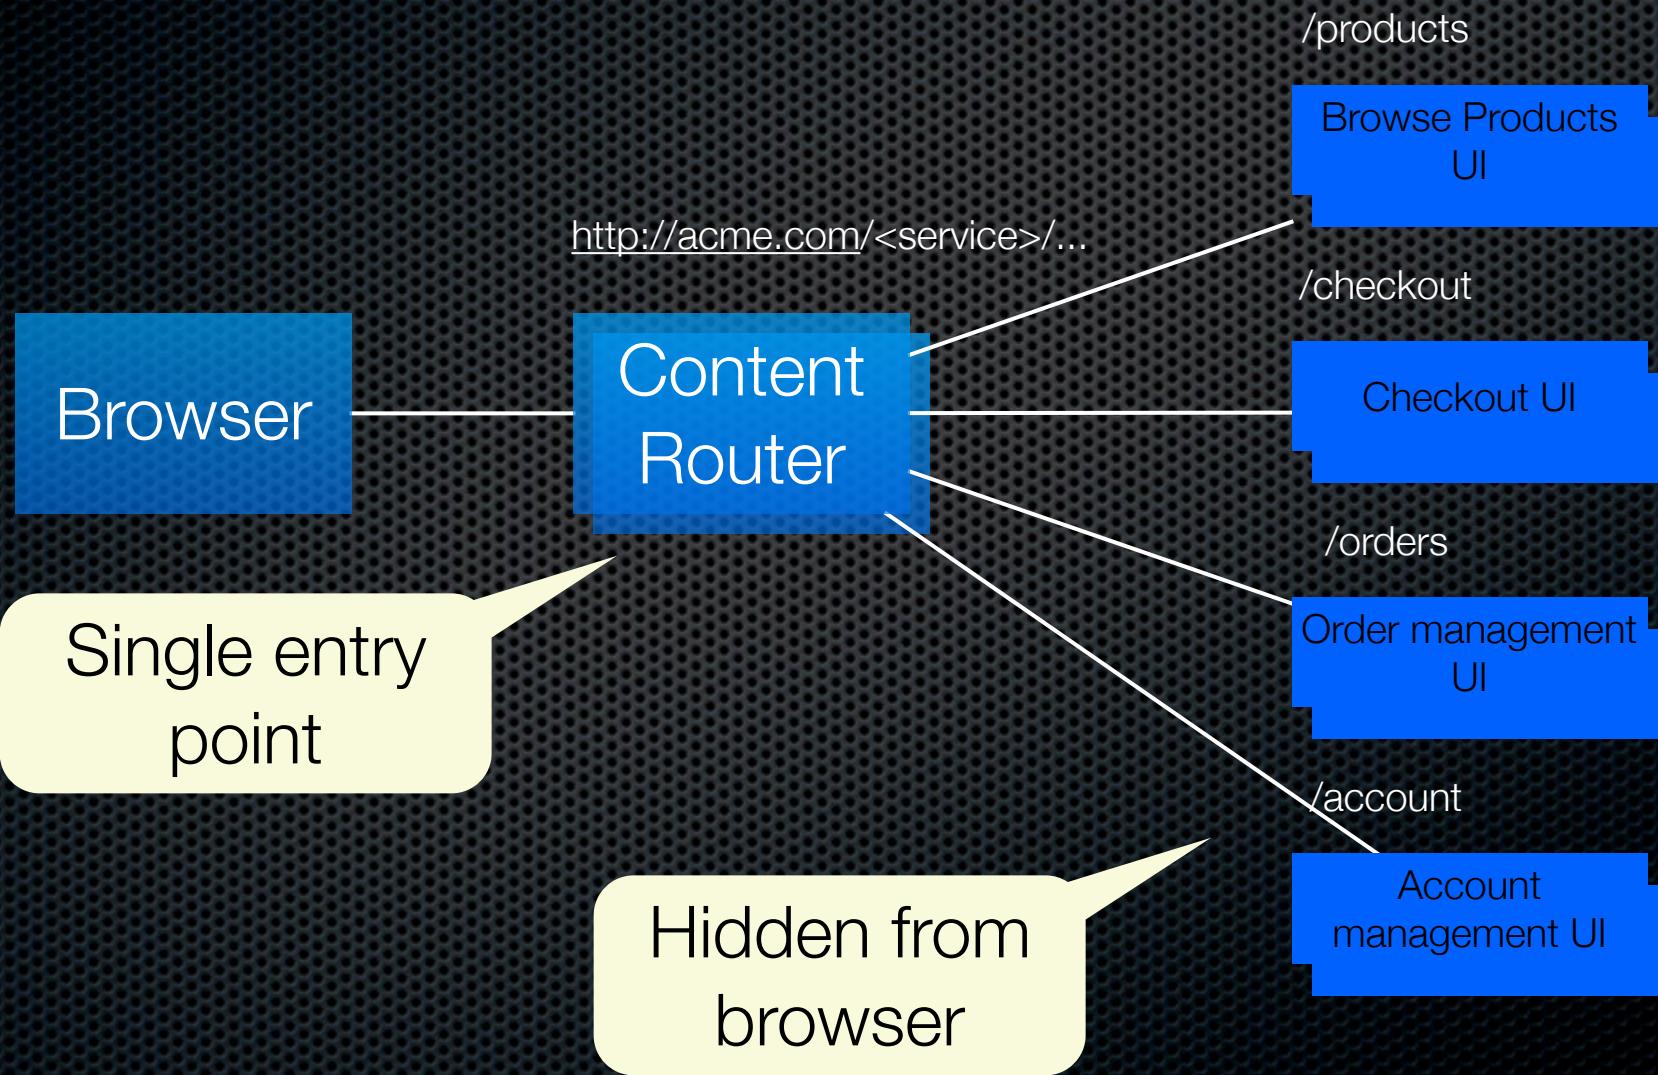
- ❖ JVM:
 - ❖ Netty, Vertx
 - ❖ Netflix Hystrix
 - ❖ ...
- ❖ Other:
 - ❖ NodeJS

How does a browser
interact with the partitioned
web application?

Partitioned web app ⇒ no longer a single base URL



The solution: single entry point that routes based on URL



How do the services communicate?

Inter-service communication options

- Synchronous HTTP ⇔ asynchronous AMQP
- Formats: JSON, XML, Protocol Buffers, Thrift, ...

Asynchronous is preferred
JSON is fashionable but binary format
is more efficient

Pros and cons of messaging

Pros

- ▣ Decouples client from server
- ▣ Message broker buffers messages
- ▣ Supports a variety of communication patterns

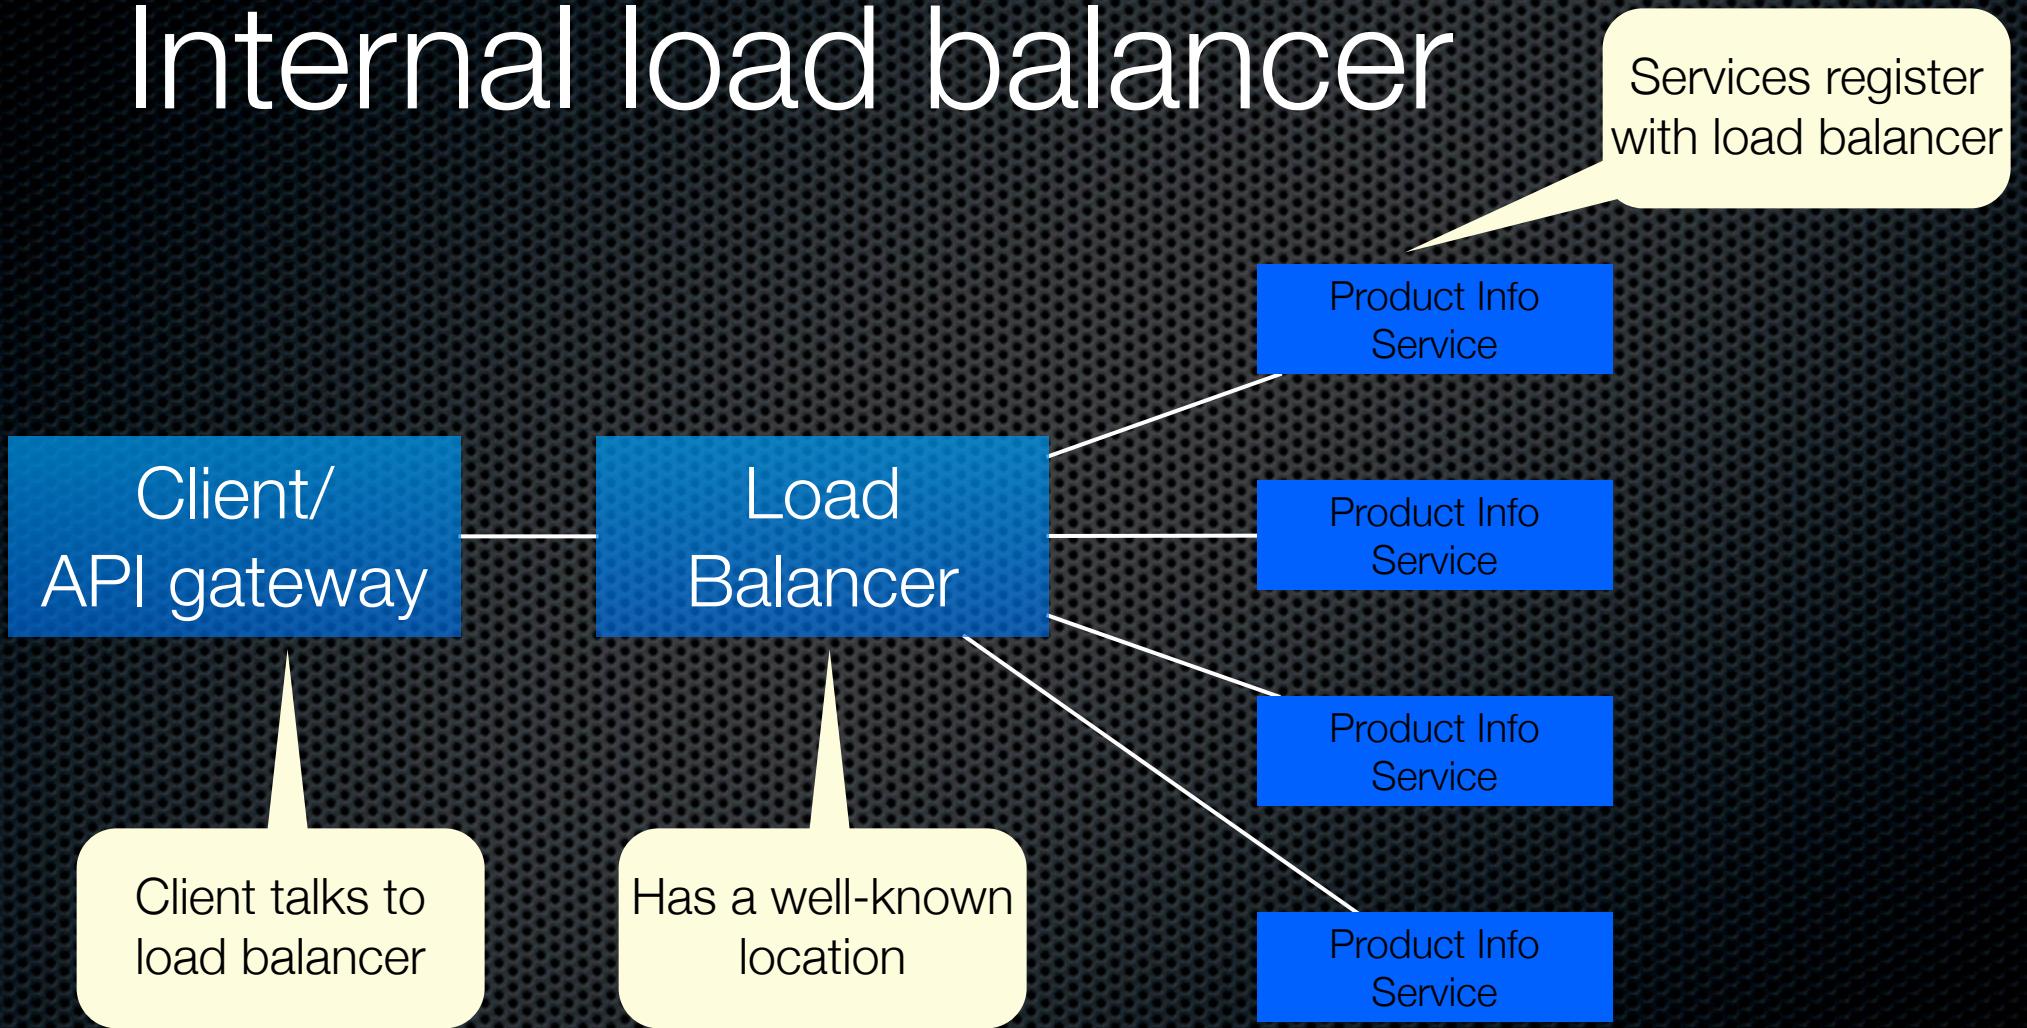
Cons

- ▣ Additional complexity of message broker
- ▣ Request/reply-style communication is more complex

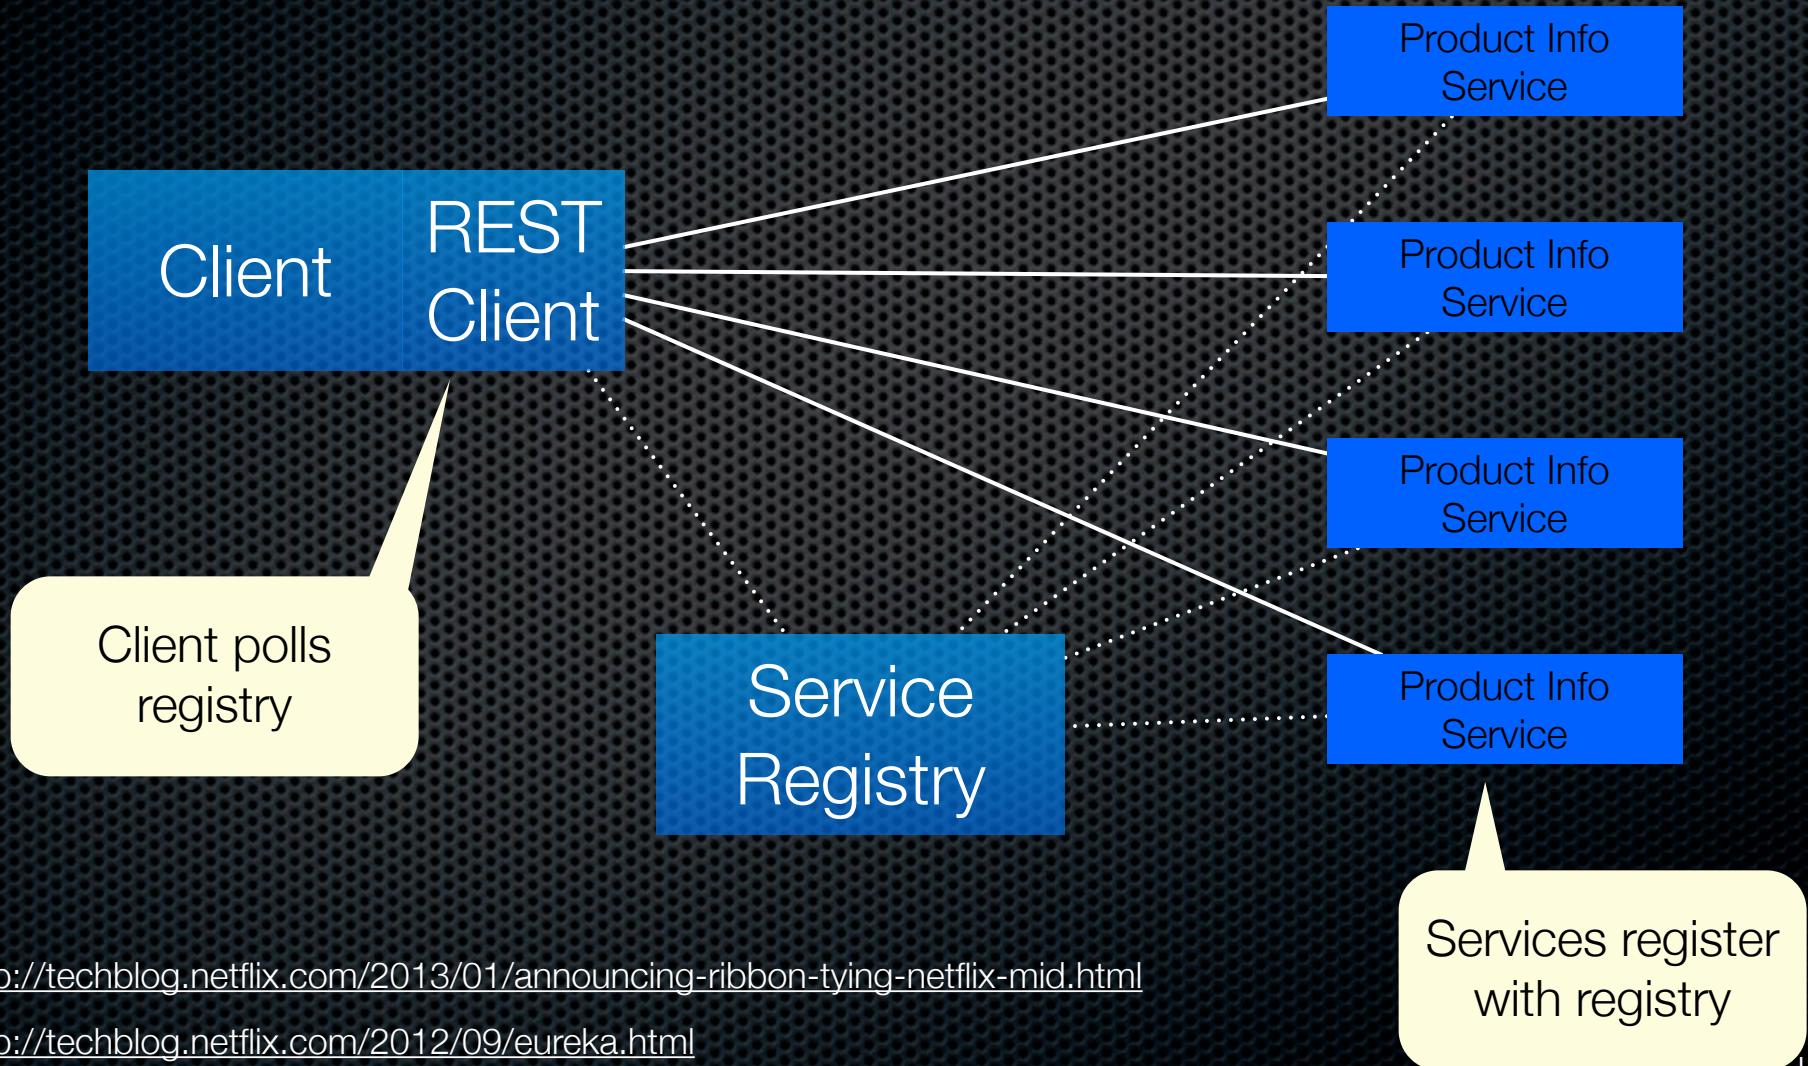
Pros and cons of HTTP

- ❖ Pros
 - ❖ Simple and familiar
 - ❖ Request/reply is easy
 - ❖ Firewall friendly
 - ❖ No intermediate broker
- ❖ Cons
 - ❖ Only supports request/reply
 - ❖ Server must be available
 - ❖ Client needs to discover URL(s) of server(s)

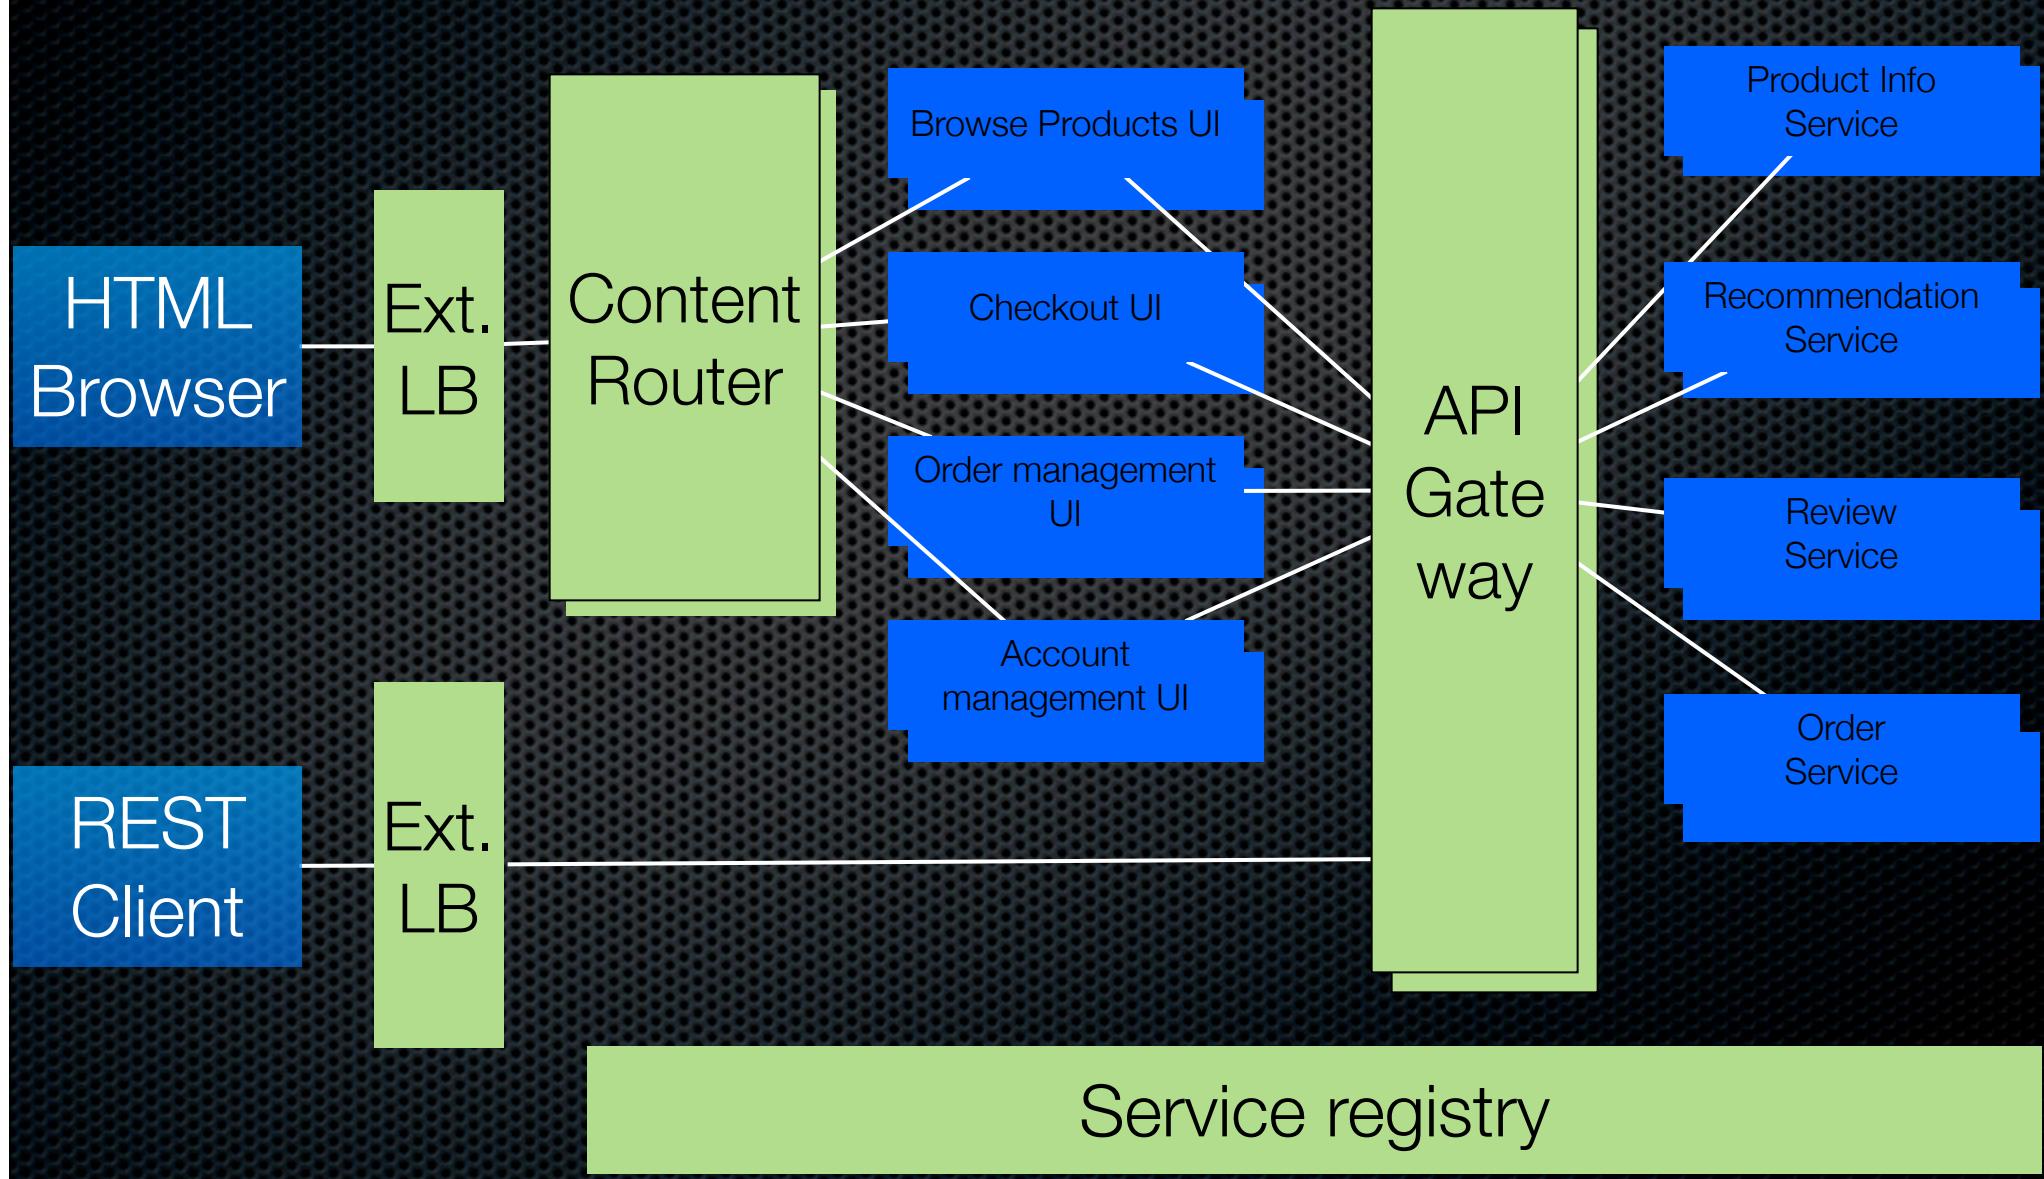
Discovery option #1: Internal load balancer



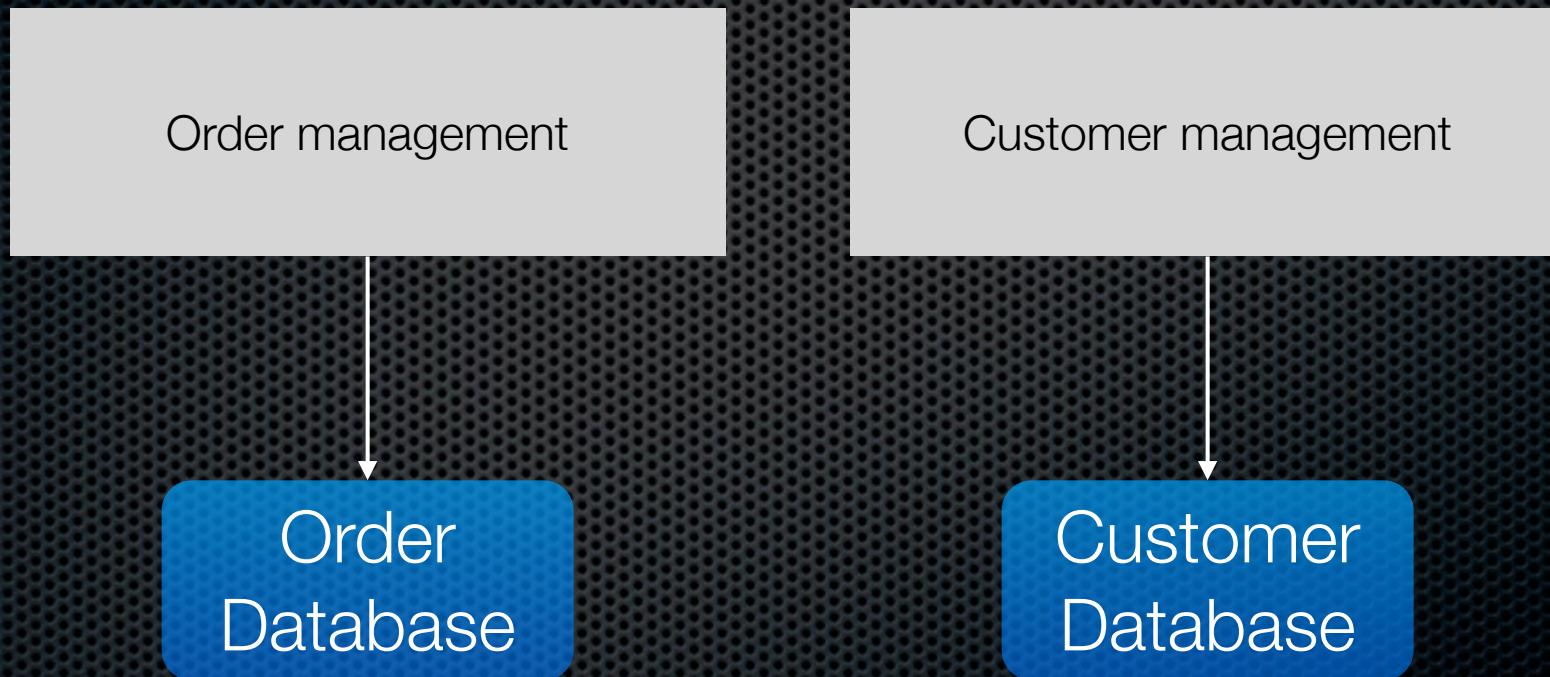
Discovery option #2: client-side load balancing



Lots of moving parts!

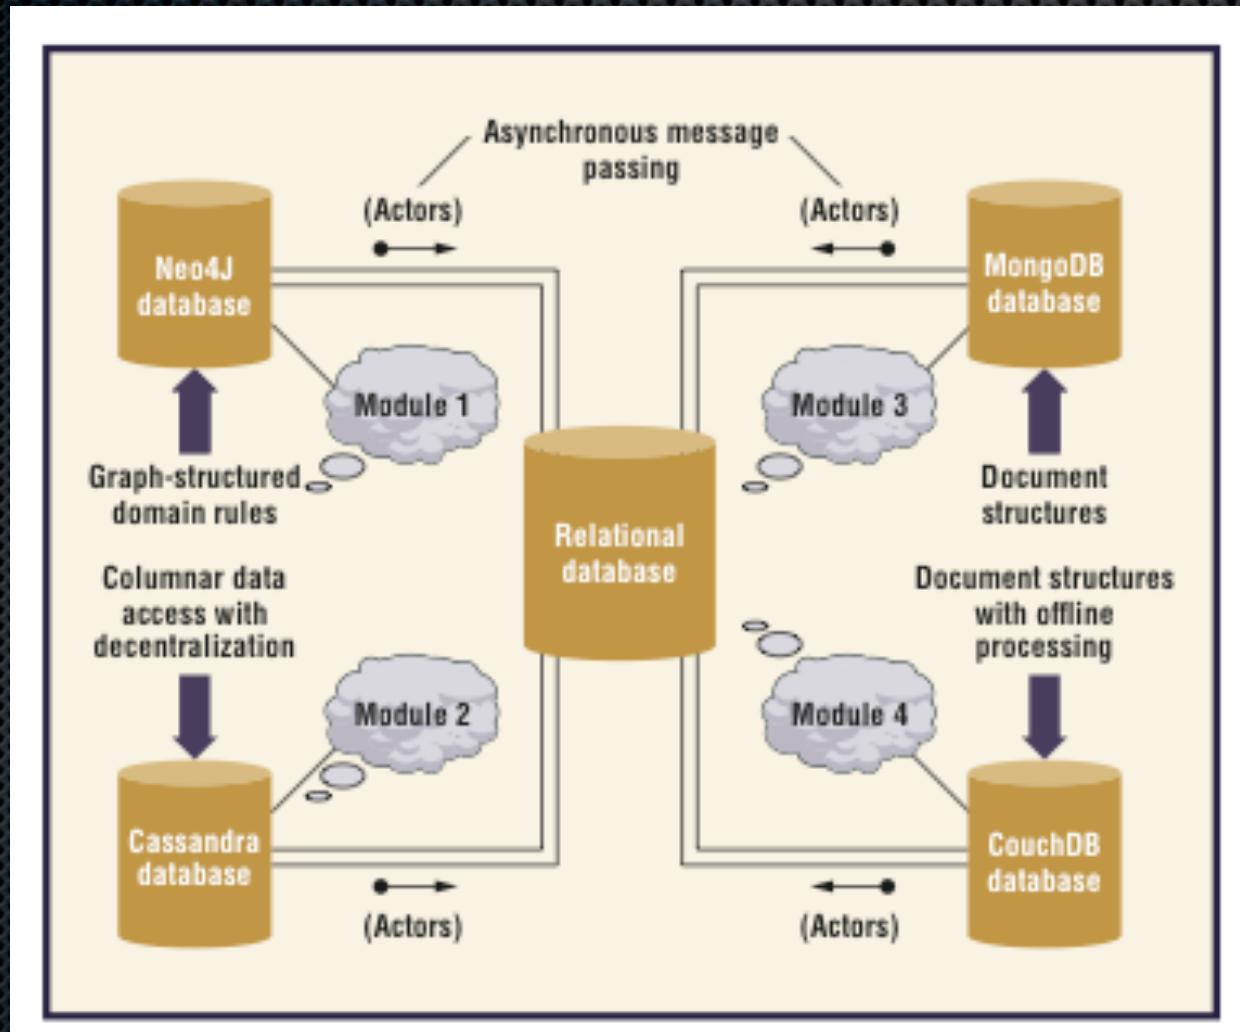


Decomposed services ⇒ decomposed databases



Separate databases ⇒ less coupling

Decomposed databases \Rightarrow polyglot persistence



Untangling orders and customers



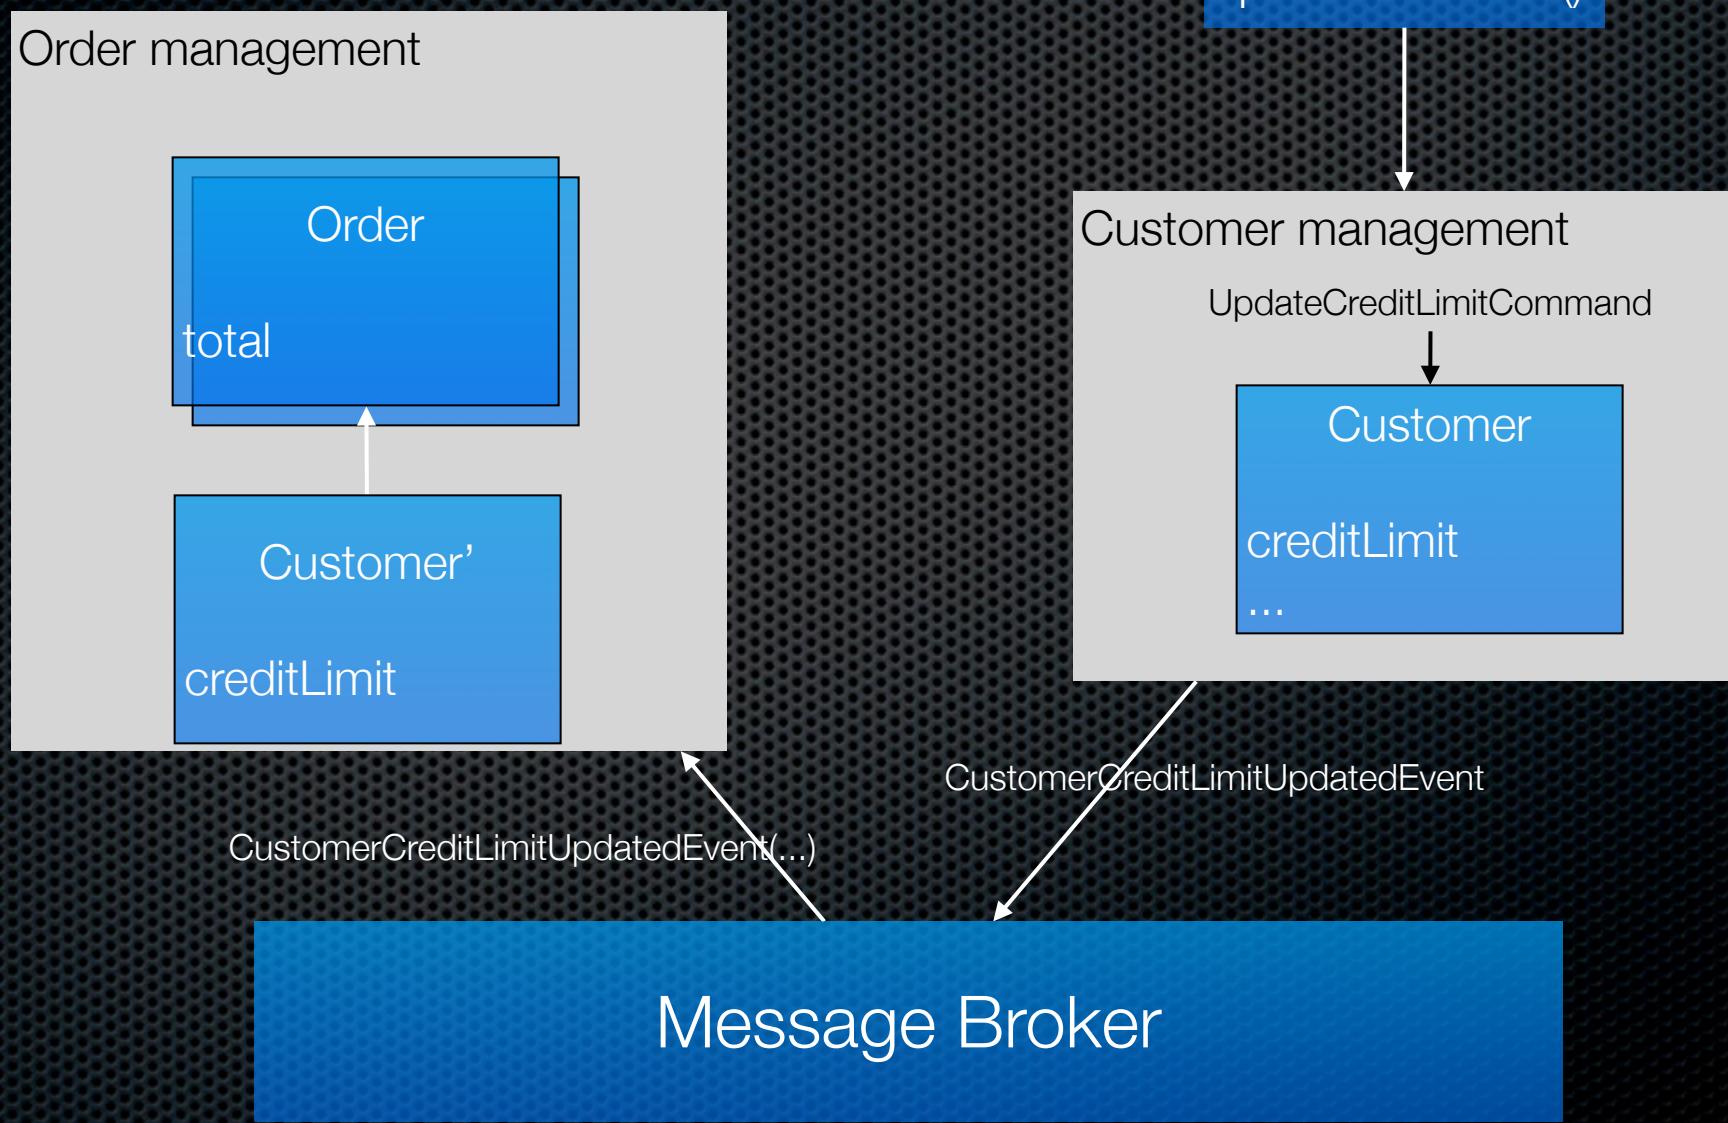
Problems

- Reads
 - Service A needs to read data owned by service B
 - e.g. OrderService needs customer's credit limit
 - e.g Customer Service needs outstanding order total
- Updates
 - “Business transaction” must update data owned by multiple services
 - e.g Creating a User => Create their ShoppingCart

You could use two-phase
commit (2PC) but ...

Use an event-driven
architecture

Event-driven architecture



Benefits and drawbacks of an event-driven architecture

- ❖ Benefits:
 - ❖ Simple infrastructure, e.g. a message broker
 - ❖ Better availability (than using 2PC)
- ❖ Drawbacks:
 - ❖ Application has to handle eventually consistent data
 - ❖ Application has to handle duplicate events

How do services publish events?

To maintain consistency the
application must
atomically publish an event
whenever
a domain object changes

How to generate events reliably?

- ❖ Database triggers
- ❖ Hibernate event listener
- ❖ Ad hoc event publishing code mixed into business logic
- ❖ Domain events - “formal” modeling of events
- ❖ Event Sourcing

Atomically publishing events

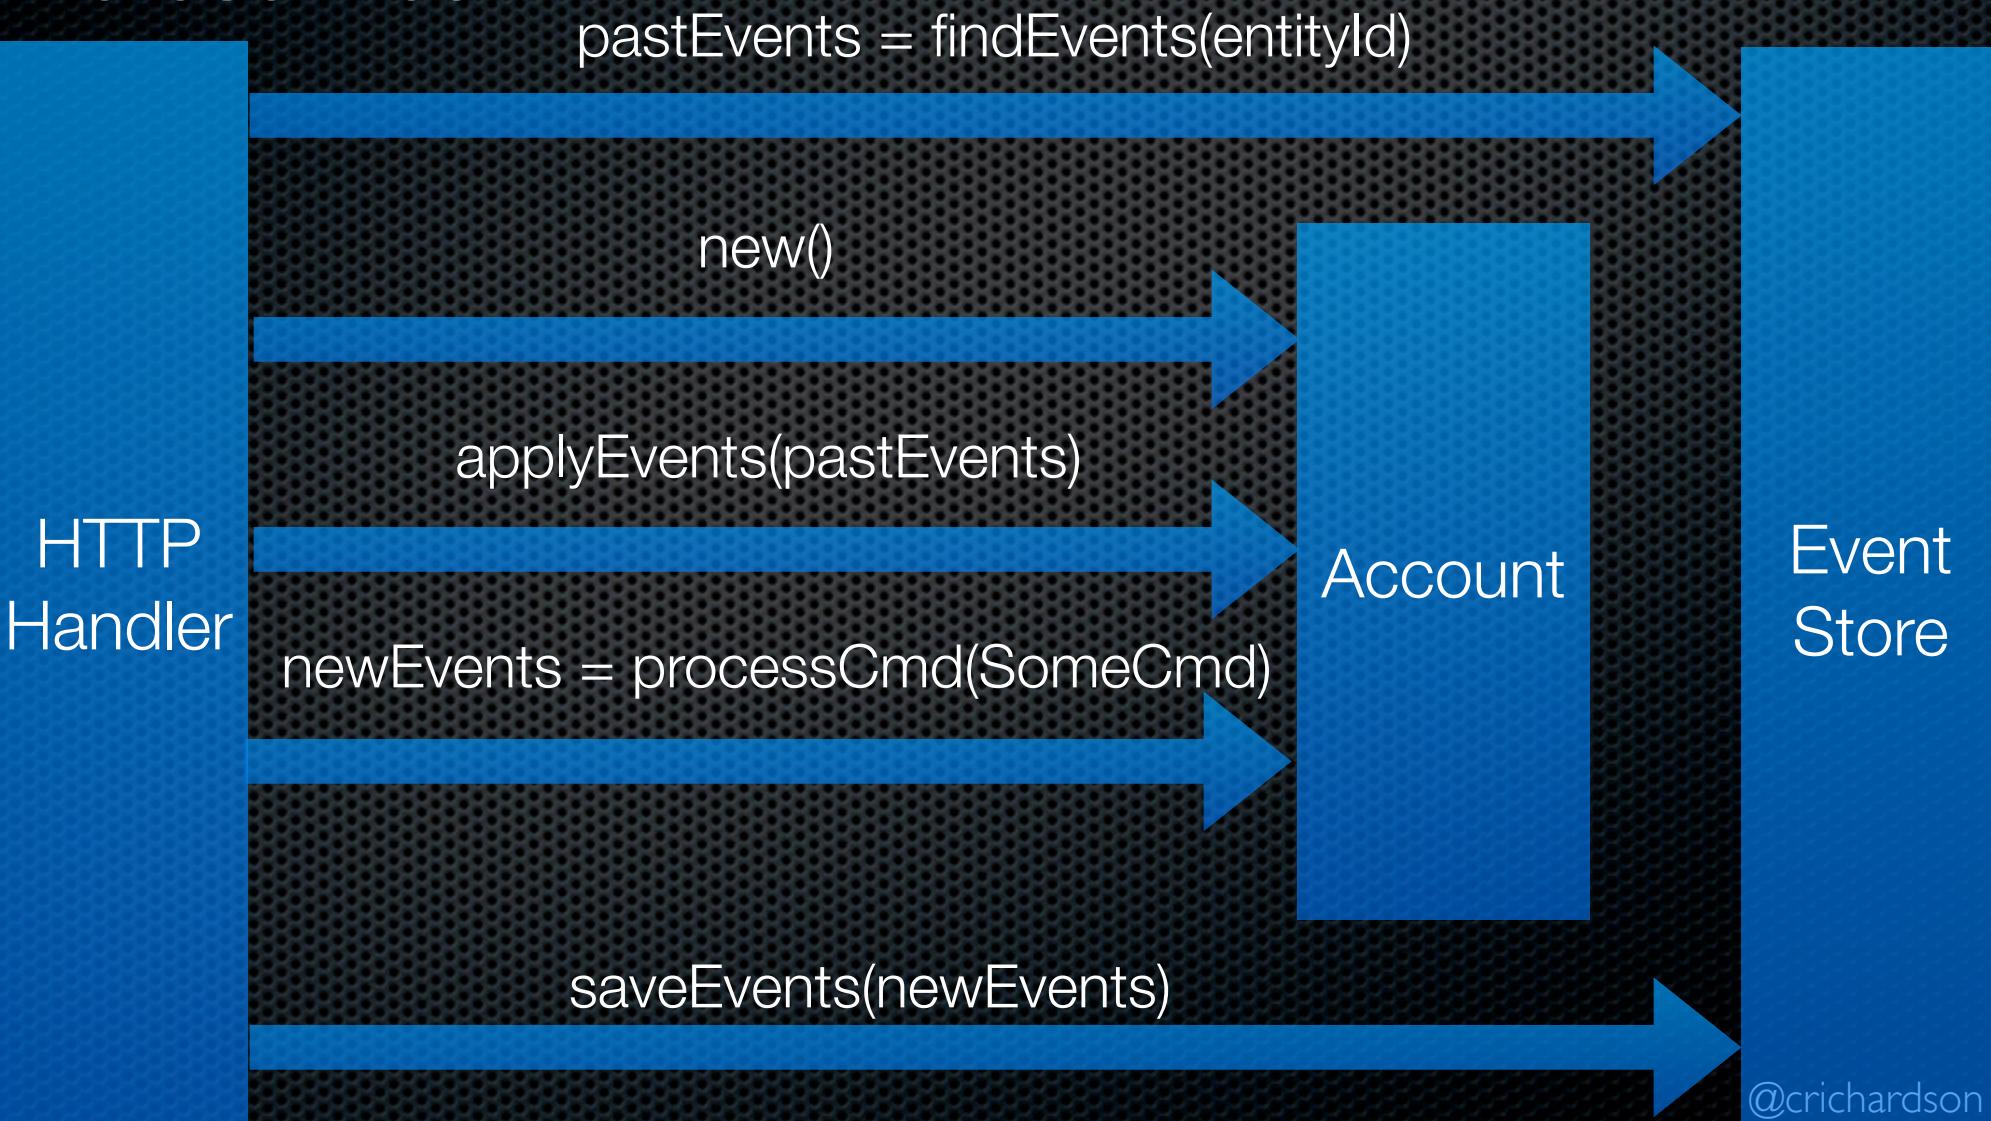
- Use distributed transactions to update database **and** publish to message
 - Database and message broker must support 2PC
 - 2PC is best avoided
- Use eventual consistency mechanism:
 1. Update database: new entity state & event to publish
 2. Publish event & mark event as published
 - Difficult to implement when using a NoSQL database :-(

Event sourcing

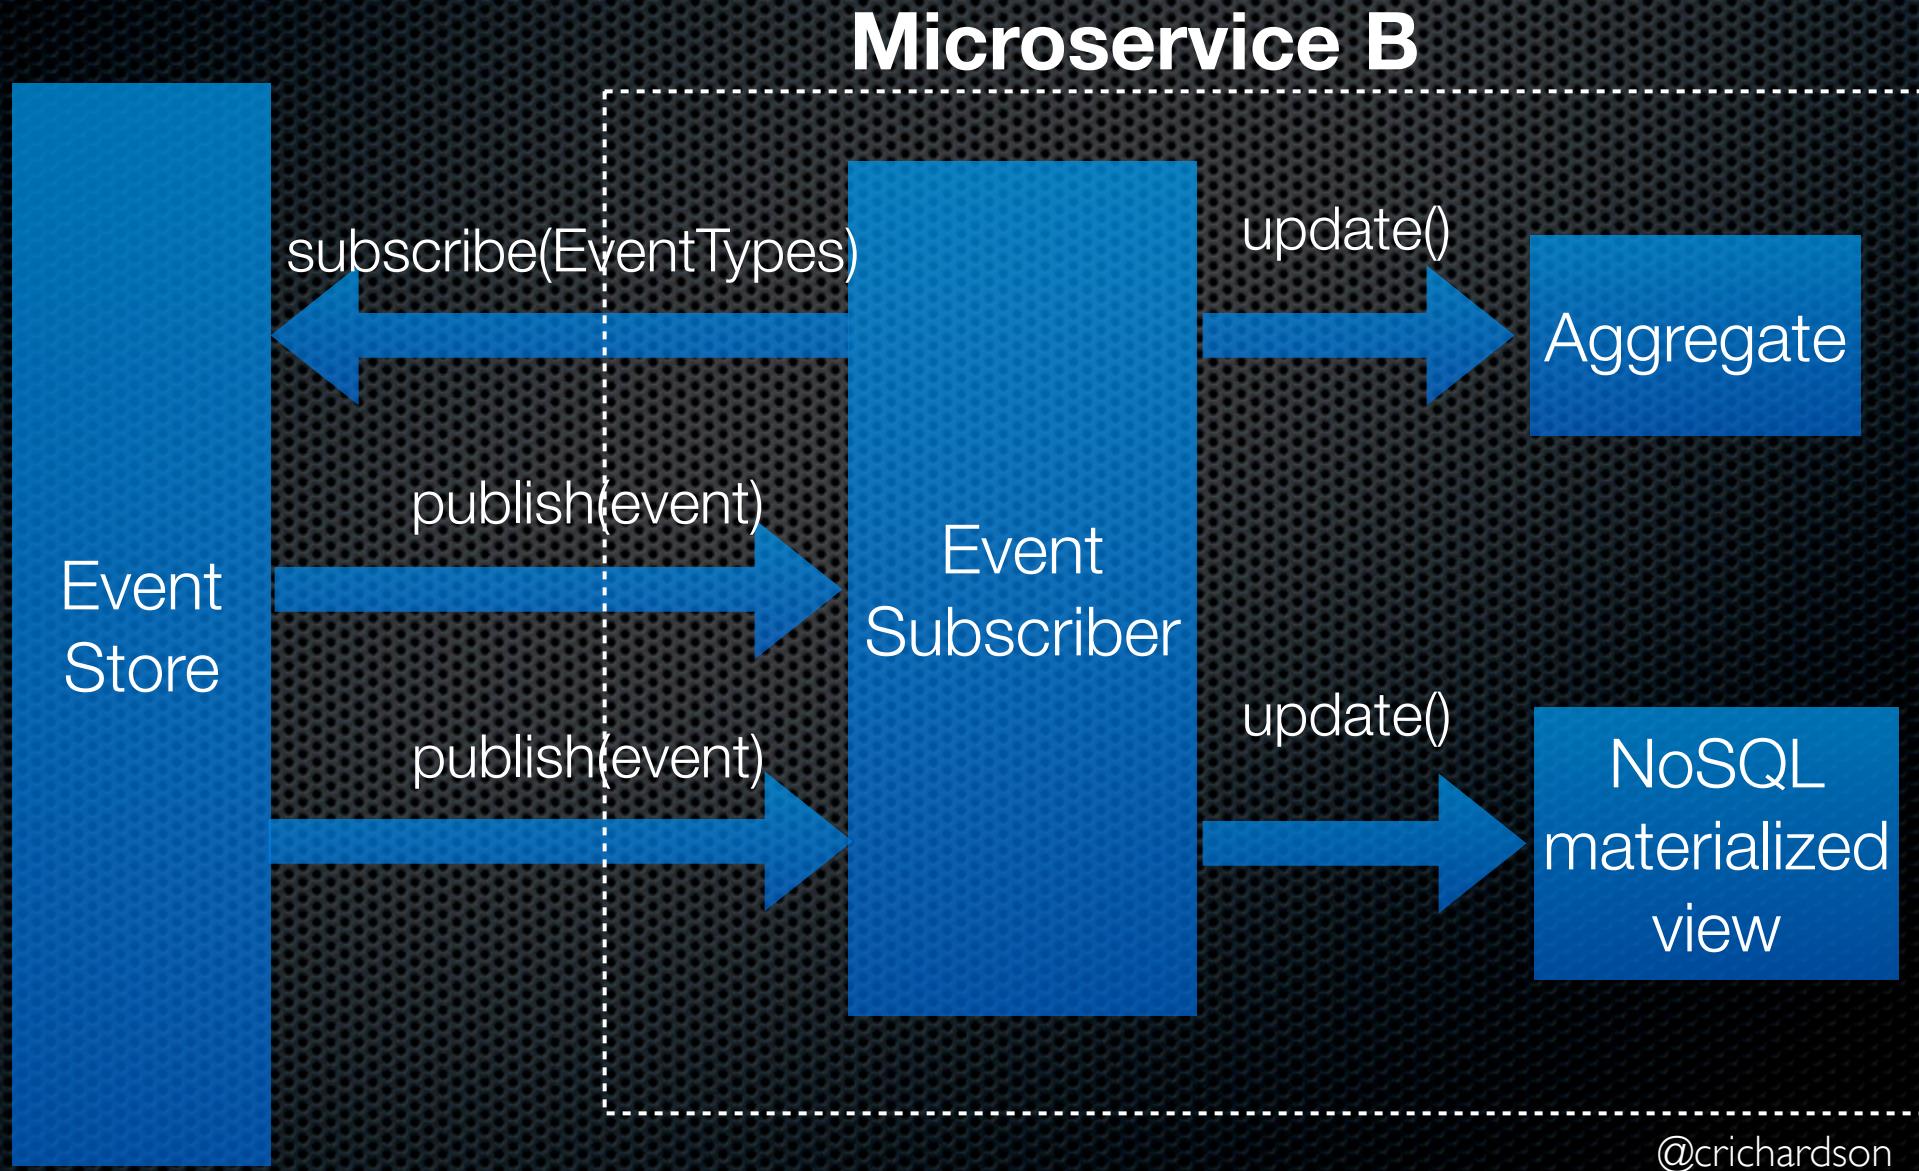
- An event-centric approach to designing domain models
 - Request becomes a command that is sent to an Aggregate
 - Aggregates handle commands by generating events
 - Apply events to an aggregate to update state
- Persist events **NOT** state
 - Replay events to recreate the current state of an aggregate
- Event Store ≈ database + message broker

Request handling in an event-sourced application

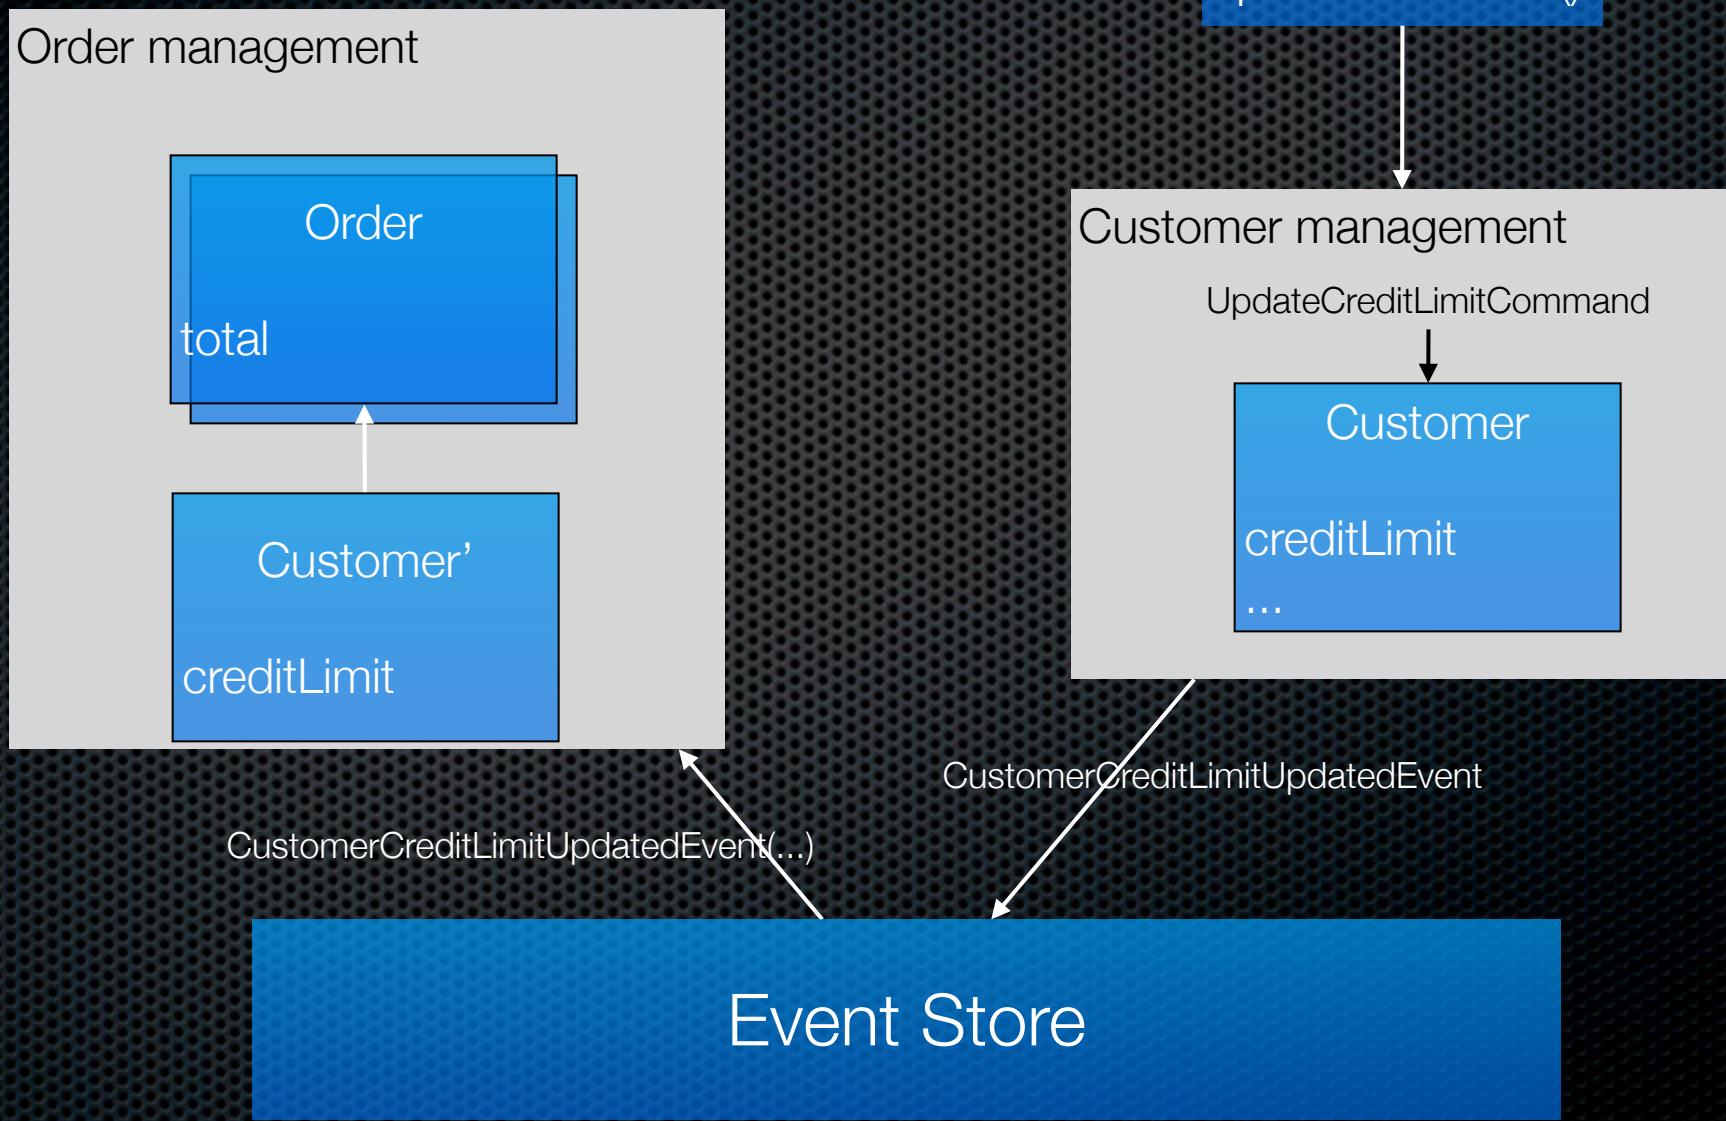
Microservice A



Event Store publishes events - consumed by other services



Using event sourcing



Unfamiliar but it solves many problems

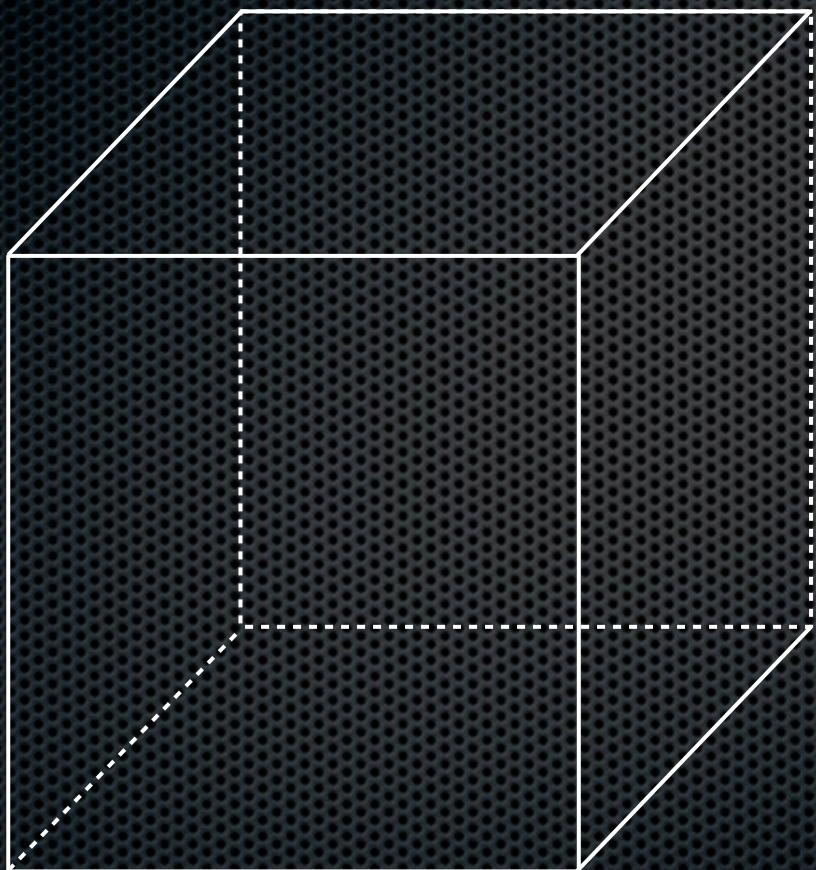
- ❖ Eliminates O/R mapping problem
- ❖ Supports both SQL and NoSQL databases
- ❖ Publishes events reliably
- ❖ Reliable eventual consistency framework
- ❖ ...

Summary

Monolithic applications are simple to develop and deploy

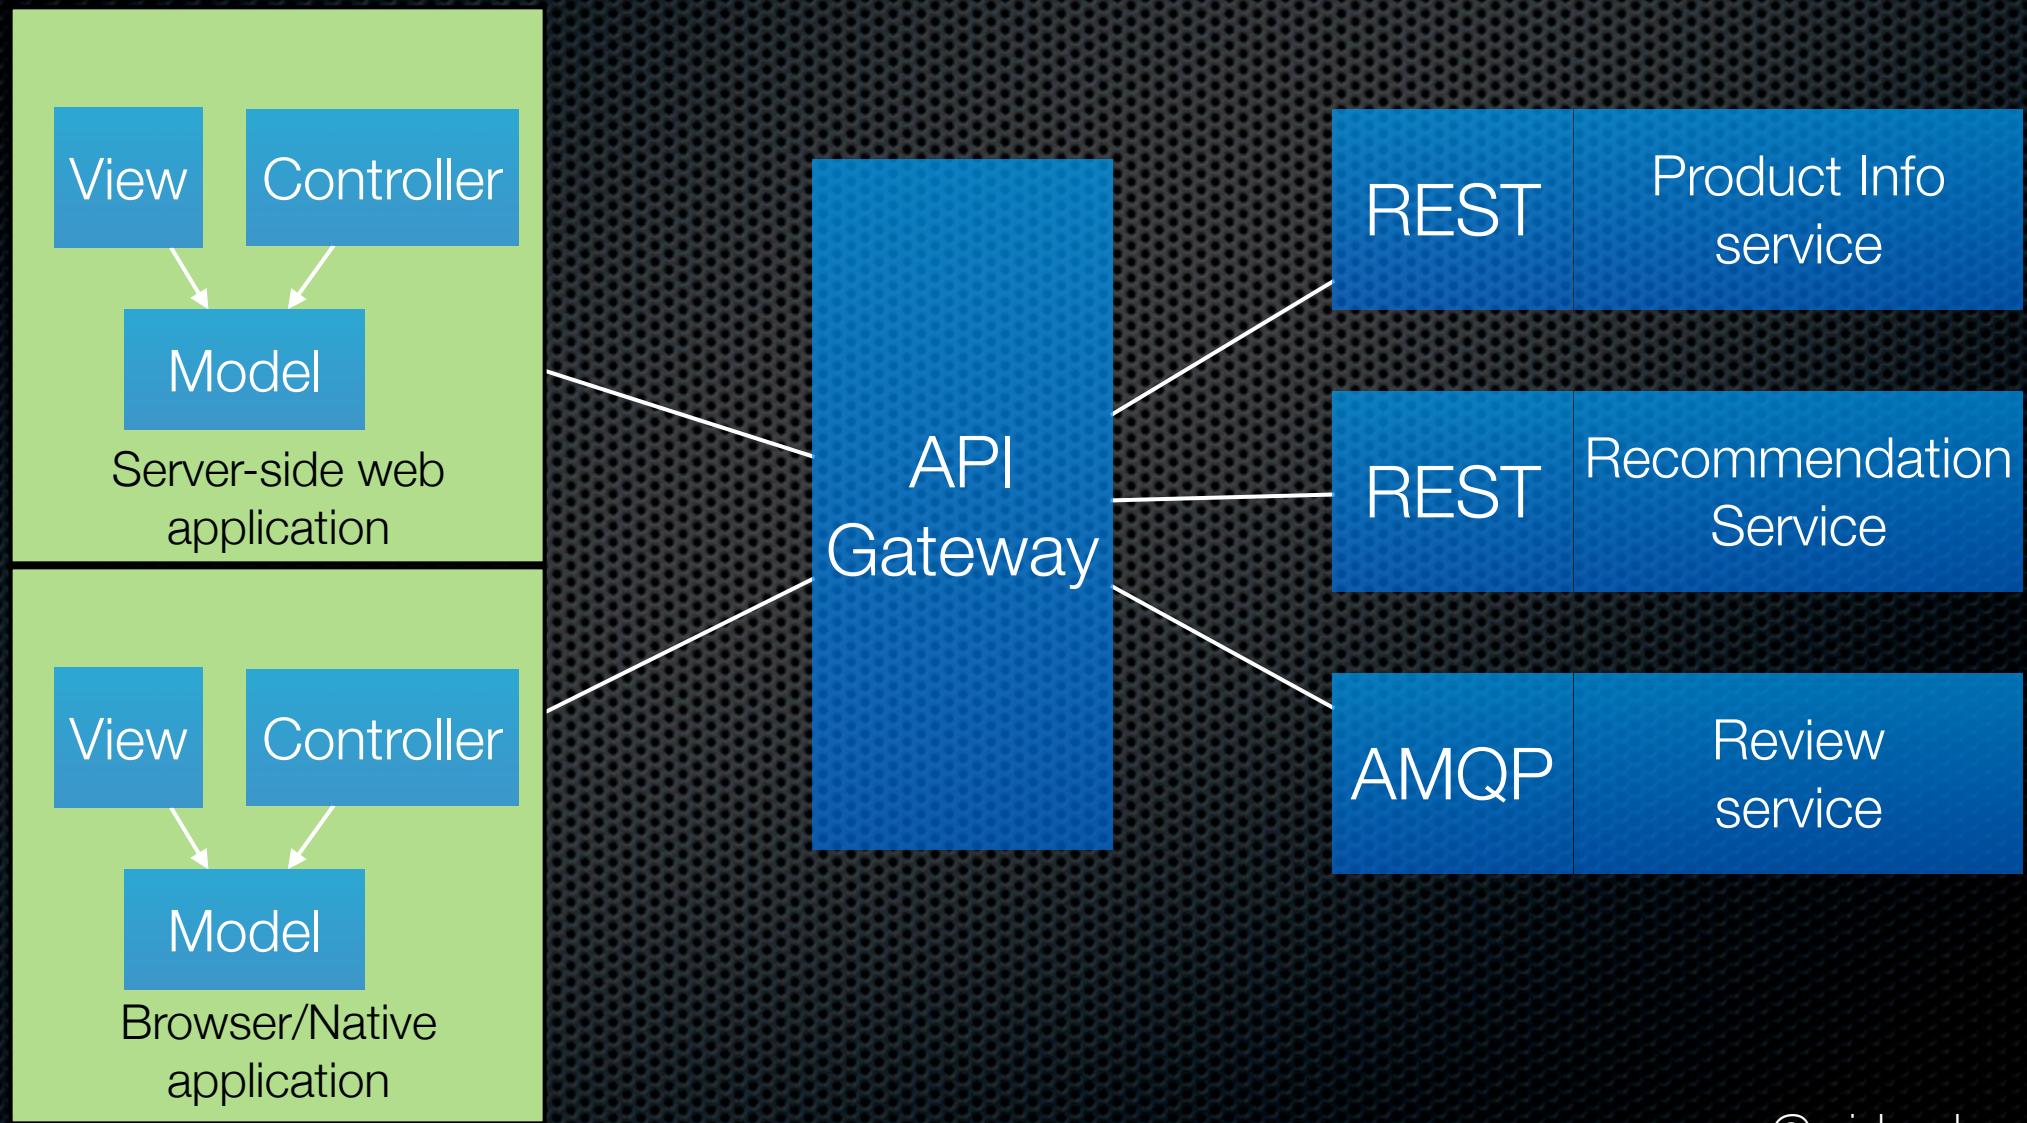
BUT have significant drawbacks

Apply the scale cube



- Modular, polyglot, and scalable applications
- Services developed, deployed and scaled independently

Use a modular, polyglot architecture



Use an event-driven
architecture

Start refactoring your monolith

